

Package ‘pomp’

August 13, 2015

Type Package

Title Statistical Inference for Partially Observed Markov Processes

Version 0.78-2

Date 2015-08-09

URL <http://kingaa.github.io/pomp>

Description Tools for working with partially observed Markov processes (POMPs, AKA stochastic dynamical systems, state-space models). 'pomp' provides facilities for implementing POMP models, simulating them, and fitting them to time series data by a variety of frequentist and Bayesian methods. It is also a platform for the implementation of new inference methods.

Depends R(>= 3.0.0), methods

Imports stats, graphics, digest, mvtnorm, deSolve, coda, subplex, nloptr

Suggests magrittr, plyr, reshape2, ggplot2, knitr

SystemRequirements For Windows users, Rtools (see <http://cran.r-project.org/bin/windows/Rtools/>).

License GPL(>= 2)

LazyData true

MailingList

Subscribe to pomp-announce@r-forge.r-project.org for announcements by going to <http://lists.r-forge.r-project.org/mailman/listinfo/pomp-announce>.

Collate aaa.R authors.R bake.R generics.R eulermultinom.R
csnippet.R pomp-fun.R plugins.R builder.R
parmat.R logmeanexp.R slice-design.R
profile-design.R sobol.R bsplines.R sannbox.R
pomp-class.R load.R pomp.R pomp-methods.R
rmeasure-pomp.R rprocess-pomp.R init-state-pomp.R
dmeasure-pomp.R dprocess-pomp.R skeleton-pomp.R
dprior-pomp.R rprior-pomp.R
simulate-pomp.R trajectory-pomp.R plot-pomp.R
pfilter.R pfilter-methods.R minim.R traj-match.R
bsmc.R bsmc2.R
mif.R mif-methods.R mif2.R mif2-methods.R
proposals.R pmcmc.R pmcmc-methods.R
nlf-funcs.R nlf-guts.R nlf-objfun.R nlf.R

probe.R probe-match.R basic-probes.R spect.R spect-match.R
 abc.R abc-methods.R covmat.R
 example.R

R topics documented:

pomp-package	3
Approximate Bayesian computation	4
B-splines	7
Bayesian sequential Monte Carlo	8
blowflies	10
Childhood disease incidence data	11
Csnippet	12
dacca	16
design	17
eulermultinom	18
Example pomp models	20
gompertz	22
Iterated filtering	22
Iterated filtering 2	27
logmeanexp	30
Low-level-interface	31
MCMC proposal distributions	35
Nonlinear forecasting	36
ou2	38
parmat	39
Particle filter	40
Particle Markov Chain Monte Carlo	43
pomp constructor	46
pomp methods	54
POMP simulation	57
Power spectrum computation and matching	58
Probe functions	61
Probes and synthetic likelihood	64
Process model plugins	67
ricker	71
rw2	71
Simulated annealing	72
sir	73
Trajectory matching	74
Utilities for reproducibility	76
Index	79

Description

The **pomp** package provides facilities for inference on time series data using partially-observed Markov process (POMP) models. These models are also known as state-space models or nonlinear stochastic dynamical systems. One can use **pomp** to fit nonlinear, non-Gaussian dynamic models to time-series data. The package is both a set of tools for data analysis and a platform upon which statistical inference methods for POMP models can be implemented.

Data analysis using pomp

The first step in using **pomp** is to encode one's model(s) and data in objects of class `pomp`. One does this via a call to `pomp`, which involves specifying the unobserved state process and the measurement process of the model. Details on this are given in the documentation for the `pomp` constructor function. Examples are given in the tutorials on the [package website](#), in the demos (`demo(package=pomp)`), and via the `pompExample` function.

pomp version 0.78-2 provides algorithms for

1. simulation of stochastic dynamical systems; see [simulate](#)
2. particle filtering (AKA sequential Monte Carlo or sequential importance sampling); see [pfilter](#)
3. the iterated filtering methods of Ionides et al. (2006, 2011, 2015); see [mif](#) and [mif2](#)
4. the nonlinear forecasting algorithm of Kendall et al. (2005); see [nlf](#)
5. the particle MCMC approach of Andrieu et al. (2010); see [pmcmc](#)
6. the probe-matching method of Kendall et al. (1999, 2005); see [probe.match](#)
7. a spectral probe-matching method (Reuman et al. 2006, 2008); see [spect.match](#)
8. synthetic likelihood a la Wood (2010); see [probe](#)
9. approximate Bayesian computation (Toni et al. 2009); see [abc](#)
10. the approximate Bayesian sequential Monte Carlo scheme of Liu & West (2001); see [bsmc](#)
11. simple trajectory matching; see [traj.match](#).

The package also provides various tools for plotting and extracting information on models and data.

Developing inference tools on the pomp platform

pomp provides a very general interface to the components of POMP models. All the inference algorithms in **pomp** interact with the models and data via this interface. One goal of the **pomp** project has been to facilitate the development of new algorithms in an environment where they can be tested and compared on a growing body of models and datasets.

The low-level interface relevant to developers is documented [here](#).

Comments, bug reports, feature requests

Contributions are welcome, as are comments, feature requests, and bug reports. See the package website <http://kingaa.github.io/pomp> for more information, access to the package new RSS feed, links to the authors' websites, references to the literature, and up-to-date versions of the package source and documentation.

Documentation

A number of tutorials, demonstrating the construction of `pomp` objects and the application of various inference algorithms, are available on the package homepage: <http://kingaa.github.io/pomp>. Several examples of the construction of `pomp` objects are provided with the package and are documented in the help pages: to view a full list of these, execute `pompExample()`. In addition, there are a number of demos, which can be viewed by executing `demo(package="pomp")`.

History

Much of the groundwork for **pomp** was laid by a working group of the National Center for Ecological Analysis and Synthesis (NCEAS), "Inference for Mechanistic Models".

License

pomp is provided under the GNU Public License (GPL).

Author(s)

Aaron A. King

References

See the package website, <http://kingaa.github.io/pomp>, for the references.

See Also

[pomp](#), [pomp low-level interface](#), [pfilter](#), [simulate](#), [mif](#), [nlf](#), [probe](#), [traj.match](#), [bsmc2](#), [pmcmc](#)

Approximate Bayesian computation

Estimation by approximate Bayesian computation (ABC)

Description

The approximate Bayesian computation (ABC) algorithm for estimating the parameters of a partially-observed Markov process.

Usage

```
## S4 method for signature pomp
abc(object, Nabc = 1, start,
     proposal, probes, scale, epsilon,
     verbose = getOption("verbose"), ...)
## S4 method for signature probed.pomp
abc(object, probes,
     verbose = getOption("verbose"), ...)
## S4 method for signature abc
abc(object, Nabc, start, proposal,
     probes, scale, epsilon,
     verbose = getOption("verbose"), ...)
## S4 method for signature abc
continue(object, Nabc = 1, ...)
## S4 method for signature abc
conv.rec(object, pars, ...)
## S4 method for signature abcList
conv.rec(object, ...)
## S4 method for signature abc
plot(x, y, pars, scatter = FALSE, ...)
## S4 method for signature abcList
plot(x, y, ...)
```

Arguments

object	An object of class pomp.
Nabc	The number of ABC iterations to perform.
start	named numeric vector; the starting guess of the parameters.
proposal	optional function that draws from the proposal distribution. Currently, the proposal distribution must be symmetric for proper inference: it is the user's responsibility to ensure that it is. Several functions that construct appropriate proposal function are provided: see MCMC proposal functions for more information.
probes	List of probes (AKA summary statistics). See probe for details.
scale	named numeric vector of scales.
epsilon	ABC tolerance.
verbose	logical; if TRUE, print progress reports.
pars	Names of parameters.
scatter	optional logical; If TRUE, draw scatterplots. If FALSE, draw traceplots.
x	abc object.
y	Ignored.
...	Additional arguments. These are currently ignored.

Running ABC

abc returns an object of class abc. One or more abc objects can be joined to form an abcList object.

Re-running ABC iterations

To re-run a sequence of ABC iterations, one can use the `abc` method on a `abc` object. By default, the same parameters used for the original ABC run are re-used (except for `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing ABC iterations

One can continue a series of ABC iterations from where one left off using the `continue` method. A call to `abc` to perform $N_{abc}=m$ iterations followed by a call to `continue` to perform $N_{abc}=n$ iterations will produce precisely the same effect as a single call to `abc` to perform $N_{abc}=m+n$ iterations. By default, all the algorithmic parameters are the same as used in the original call to `abc`. Additional arguments will override the defaults.

Methods

Methods that can be used to manipulate, display, or extract information from an `abc` object:

`conv.rec(object, pars)` returns the columns of the convergence-record matrix corresponding to the names in `pars`. By default, all rows are returned.

`c` Concatenates `abc` objects into an `abcList`.

`plot` Diagnostic plots.

`covmat(object, start, thin, expand)` computes the empirical covariance matrix of the ABC samples beginning with iteration `start` and thinning by factor `thin`. It expands this by a factor expand^2/n , where n is the number of parameters estimated. The intention is that the resulting matrix is a suitable input to the proposal function `mvn.rw`.

Author(s)

Edward L. Ionides, Aaron A. King

References

T. Toni and M. P. H. Stumpf, Simulation-based model selection for dynamical systems in systems and population biology, *Bioinformatics* 26:104–110, 2010.

T. Toni, D. Welch, N. Strelkowa, A. Ipsen, and M. P. H. Stumpf, Approximate Bayesian computation scheme for parameter inference and model selection in dynamical systems *Journal of the Royal Society, Interface* 6:187–202, 2009.

See Also

[pomp](#), [probe](#), [MCMC proposal distributions](#), and the tutorials on the [package website](#).

B-splines

*B-spline bases***Description**

These functions generate B-spline basis functions. `bspline.basis` gives a basis of spline functions. `periodic.bspline.basis` gives a basis of periodic spline functions.

Usage

```
bspline.basis(x, nbasis, degree = 3, names = NULL)
periodic.bspline.basis(x, nbasis, degree = 3, period = 1, names = NULL)
```

Arguments

<code>x</code>	Vector at which the spline functions are to be evaluated.
<code>nbasis</code>	The number of basis functions to return.
<code>degree</code>	Degree of requested B-splines.
<code>period</code>	The period of the requested periodic B-splines.
<code>names</code>	optional; the names to be given to the basis functions. These will be the column-names of the matrix returned. If the names are specified as a format string (e.g., "basis%d"), <code>sprintf</code> will be used to generate the names from the column number. If a single non-format string is specified, the names will be generated by <code>paste</code> -ing name to the column number. One can also specify each column name explicitly by giving a length- <code>nbasis</code> string vector. By default, no column-names are given.

Value

<code>bspline.basis</code>	Returns a matrix with <code>length(x)</code> rows and <code>nbasis</code> columns. Each column contains the values one of the spline basis functions.
<code>periodic.bspline.basis</code>	Returns a matrix with <code>length(x)</code> rows and <code>nbasis</code> columns. The basis functions returned are periodic with period <code>period</code> .

C API

Access to the underlying C routines is available: see the header file 'pomp.h' for definition and documentation of the C API. At an R prompt, execute

```
file.show(system.file("include/pomp.h", package="pomp"))
```

to view this file.

Author(s)

Aaron A. King

Examples

```
x <- seq(0,2,by=0.01)
y <- bspline.basis(x,degree=3,nbasis=9,names="basis")
matplot(x,y,type=l,ylim=c(0,1.1))
lines(x,apply(y,1,sum),lwd=2)

x <- seq(-1,2,by=0.01)
y <- periodic.bspline.basis(x,nbasis=5,names="spline%d")
matplot(x,y,type=l)
```

Bayesian sequential Monte Carlo

The Liu and West Bayesian particle filter

Description

Modified versions of the Liu and West (2001) algorithm.

Usage

```
## S4 method for signature pomp
bsmc2(object, params, Np, est, smooth = 0.1,
      tol = 1e-17, verbose = getOption("verbose"), max.fail = 0,
      transform = FALSE, ...)
## S4 method for signature pomp
bsmc(object, params, Np, est, smooth = 0.1,
     ntries = 1, tol = 1e-17, lower = -Inf, upper = Inf,
     verbose = getOption("verbose"), max.fail = 0,
     transform = FALSE, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> or inheriting class <code>pomp</code> .
<code>params</code> , <code>Np</code>	Specifications for the prior distribution of particles. See details below.
<code>est</code>	Names of the rows of <code>params</code> that are to be estimated. No updates will be made to the other parameters. If <code>est</code> is not specified, all parameters for which there is variation in <code>params</code> will be estimated.
<code>smooth</code>	Kernel density smoothing parameters. The compensating shrinkage factor will be $\sqrt{1-\text{smooth}^2}$. Thus, <code>smooth=0</code> means that no noise will be added to parameters. Generally, the value of <code>smooth</code> should be chosen close to 0 (i.e., $\text{shrink} \sim 0.1$).
<code>ntries</code>	Number of draws from <code>rprocess</code> per particle used to estimate the expected value of the state process at time $t+1$ given the state and parameters at time t .
<code>tol</code>	Particles with log likelihood below <code>tol</code> are considered to be “lost”. A filtering failure occurs when, at some time point, all particles are lost. When all particles are lost, the conditional log likelihood at that time point is set to be $\log(\text{tol})$.

lower, upper	optional; lower and upper bounds on the priors. This is useful in case there are box constraints satisfied by the priors. The posterior is guaranteed to lie within these bounds.
verbose	logical; if TRUE, print diagnostic messages.
max.fail	The maximum number of filtering failures allowed. If the number of filtering failures exceeds this number, execution will terminate with an error.
transform	logical; if TRUE, the algorithm operates on the transformed scale.
...	currently ignored.

Details

There are two ways to specify the prior distribution of particles. If `params` is unspecified or is a named vector, N_p draws are made from the prior distribution, as specified by `rprior`. Alternatively, `params` can be specified as an $n_{\text{pars}} \times N_p$ matrix (with rownames).

`bsmc` uses version of the original algorithm that includes a plug-and-play auxiliary particle filter. `bsmc2` discards this auxiliary particle filter and appears to give superior performance for the same amount of effort.

Value

An object of class “`bsmcd.pomp`”. The “`params`” slot of this object will hold the parameter posterior medians. The slots of this class include:

<code>post</code>	A matrix containing draws from the approximate posterior distribution.
<code>prior</code>	A matrix containing draws from the prior distribution (identical to <code>params</code> on call).
<code>eff.sample.size</code>	A vector containing the effective number of particles at each time point.
<code>smooth</code>	The smoothing parameter used (see above).
<code>nfail</code>	The number of filtering failures encountered.
<code>cond.log.evidence</code>	A vector containing the conditional log evidence scores at each time point.
<code>log.evidence</code>	The estimated log evidence.
<code>weights</code>	The resampling weights for each particle.

Author(s)

Michael Lavine (lavine at math dot umass dot edu), Matthew Ferrari (mferrari at psu dot edu), Aaron A. King (kingaa at umich dot edu), Edward L. Ionides (ionides at umich dot edu)

References

Liu, J. and M. West. Combining Parameter and State Estimation in Simulation-Based Filtering. In A. Doucet, N. de Freitas, and N. J. Gordon, editors, *Sequential Monte Carlo Methods in Practice*, pages 197-224. Springer, New York, 2001.

See Also

[pomp](#), [pfilter](#)

blowflies

Model for Nicholson's blowflies.

Description

blowflies1 and blowflies2 are pomp objects encoding stochastic delay-difference models.

Details

The data are from "population I", a control culture in one of A. J. Nicholson's experiments with the Australian sheep-blowfly *Lucilia cuprina*. The experiment is described on pp. 163–4 of Nicholson (1957). Unlimited quantities of larval food were provided; the adult food supply (ground liver) was constant at 0.4g per day. The data were taken from the table provided by Brillinger et al. (1980).

The models are discrete delay equations:

$$R(t+1) \sim \text{Poisson}(PN(t-\tau) \exp(-N(t-\tau)/N_0)e(t+1)\Delta t)$$

$$S(t+1) \sim \text{binomial}(N(t), \exp(-\delta\epsilon(t+1)\Delta t))$$

$$N(t) = R(t) + S(t)$$

where $e(t)$ and $\epsilon(t)$ are Gamma-distributed i.i.d. random variables with mean 1 and variances $\sigma_p^2/\Delta t$, $\sigma_d^2/\Delta t$, respectively. blowflies1 has a timestep (Δt) of 1 day, and blowflies2 has a timestep of 2 days. The process model in blowflies1 thus corresponds exactly to that studied by Wood (2010). The measurement model in both cases is taken to be

$$y(t) \sim \text{negbin}(N(t), 1/\sigma_y^2)$$

, i.e., the observations are assumed to be negative-binomially distributed with mean $N(t)$ and variance $N(t) + (\sigma_y N(t))^2$.

Do

```
file.show(system.file("examples", "blowflies.R", package="pomp"))
```

to view the code that constructs these pomp objects.

References

- A. J. Nicholson (1957) The self-adjustment of populations to change. Cold Spring Harbor Symposia on Quantitative Biology, **22**, 153–173.
- Y. Xia and H. Tong (2011) Feature Matching in Time Series Modeling. *Statistical Science* **26**, 21–46.
- E. L. Ionides (2011) Discussion of “Feature Matching in Time Series Modeling” by Y. Xia and H. Tong. *Statistical Science* **26**, 49–52.
- S. N. Wood (2010) Statistical inference for noisy nonlinear ecological dynamic systems. *Nature* **466**, 1102–1104.

W. S. C. Gurney, S. P. Blythe, and R. M. Nisbet (1980) Nicholson's blowflies revisited. *Nature* **287**, 17–21.

D. R. Brillinger, J. Guckenheimer, P. Guttorp and G. Oster (1980) Empirical modelling of population time series: The case of age and density dependent rates. in G. Oster (ed.), *Some Questions in Mathematical Biology*, vol. 13, pp. 65–90. American Mathematical Society, Providence.

See Also

[pomp](#)

Examples

```
pompExample(blowflies)
plot(blowflies1)
plot(blowflies2)
```

Childhood disease incidence data

Historical childhood disease incidence data

Description

LondonYorke is a data frame containing the monthly number of reported cases of chickenpox, measles, and mumps from two American cities (Baltimore and New York) in the mid-20th century (1928–1972).

ewmeas and ewcitmeas are data frames containing weekly reported cases of measles in England and Wales. ewmeas records the total measles reports for the whole country, 1948–1966. One questionable data point has been replaced with an NA. ewcitmeas records the incidence in seven English cities 1948–1987. These data were kindly provided by Ben Bolker, who writes: “Most of these data have been manually entered from published records by various people, and are prone to errors at several levels. All data are provided as is; use at your own risk.”

Usage

```
LondonYorke
ewmeas
ewcitmeas
```

References

W. P. London and J. A. Yorke, Recurrent Outbreaks of Measles, Chickenpox and Mumps: I. Seasonal Variation in Contact Rates, *American Journal of Epidemiology*, 98:453–468, 1973.

Examples

```

plot(cases~time,data=LondonYorke,subset=disease=="measles",type=n,main="measles",bty=1)
lines(cases~time,data=LondonYorke,subset=disease=="measles"&town=="Baltimore",col="red")
lines(cases~time,data=LondonYorke,subset=disease=="measles"&town=="New York",col="blue")
legend("topright",legend=c("Baltimore","New York"),lty=1,col=c("red","blue"),bty=n)

plot(
  cases~time,
  data=LondonYorke,
  subset=disease=="chickenpox"&town=="New York",
  type=l,col="blue",main="chickenpox, New York",
  bty=1
)

plot(
  cases~time,
  data=LondonYorke,
  subset=disease=="mumps"&town=="New York",
  type=l,col="blue",main="mumps, New York",
  bty=1
)

plot(reports~time,data=ewmeas,type=l)

plot(reports~date,data=ewcitmeas,subset=city=="Liverpool",type=l)

```

Csnippet

C code snippets for accelerating computations

Description

For including snippets of C code in pomp objects.

Usage

```
Csnippet(text)
```

Arguments

text character; a snippet of C code.

Value

An object of class Csnippet.

Using C snippets to accelerate computations

From version 0.50, **pomp** provides a facility whereby users can define their model's components using inline C code. Furnishing one or more Csnippets as arguments to the **pomp** constructor causes them to be written to a C file stored in the R session's temporary directory, which is then compiled (via `R CMD SHLIB`) into a dynamically loadable shared object file. This is then loaded as needed.

Note to Windows and Mac users: By default, your R installation may not support `R CMD SHLIB`. The package website contains installation instructions that show how to enable this powerful feature of R.

In writing a Csnippet one must bear in mind both the *goal* of the snippet, i.e., what computation it is intended to perform, and the *context* in which it will be executed. Details of both of these are given below in the form of rules governing the use of Csnippets. Illustrative examples are given in the tutorials on the [package website](#).

General rules for writing C snippets

1. C snippets must be valid C. They will be embedded verbatim in a template file which will then be compiled by a call to `R CMD SHLIB`. If the resulting file does not compile, an error message will be generated. No attempt is made by **pomp** to interpret this message. Typically, compilation errors are due to either invalid C syntax or undeclared variables.
2. State variables, parameters, observables, and covariates must be left undeclared in the snippet. State variables and parameters must be declared in either the `statenames` or `paramnames` argument to **pomp**, as appropriate. Compiler errors that complain about undeclared state variables or parameters are usually due to failure to include these parameters in the appropriate vector of names.
3. A C snippet can declare local variables. Be careful not to use names that match those of state variables, observables, or parameters. The latter must never be declared within a snippet.
4. Names of observables are determined by their names in the data. They must be referred to in measurement model snippets (`rmeasure` or `dmeasure`) by those names.
5. If the **pomp** object contains a table of covariates (see [pomp](#)), then the variables in the covariate table will be available, by their names, in the context within which the snippet is executed.
6. R variables with names containing dots (`'.'`) are replaced in the C codes by variable names in which all dots have been replaced by underscores (`'_'`).
7. The header `'R.h'`, provided with **R**, will be included in the generated C file, making all of the **R C API** available for use in the snippet.
8. The header `'pomp.h'`, provided with **pomp**, will also be included, making all of the **pomp C API** available for use in every snippet. Do

```
file.show(system.file("include/pomp.h", package="pomp"))
```

to view this header file.

9. Snippets of C code passed to the `globals` argument of **pomp** will be included at the head of the generated C file. This can be used to declare global variables, define useful functions, and include arbitrary header files.

Rules for writing `rmeasure` snippets

1. The goal of such a snippet is to fill the observables with random values drawn from the measurement model distribution. Accordingly, each observable should be assigned a new value.
2. In addition to the states, parameters, covariates (if any), and observables, the variable `t`, containing the time of the observation, will be defined in the context in which the snippet is executed.

Rules for writing `dmeasure` snippets

1. The goal of such a snippet is to fill the `lik` variable with the likelihood of the data given the state. Alternatively, if `give_log=1`, `lik` should be filled with the log likelihood.
2. In addition to the states, parameters, covariates (if any), and observables, the variable `t`, containing the time of the observation, will be defined in the context in which the snippet is executed.

Rules for writing `euler.sim` and `discrete.time.sim` snippets

1. The goal of such a snippet is to replace the state variables with their new random values at the end of the time interval. Accordingly, each state variable should be over-written with its new value.
2. In addition to the states, parameters, covariates (if any), and observables, the variables `t` and `dt`, containing respectively the time at the beginning of the Euler step and the Euler step's duration, will be defined in the context in which the snippet is executed.

Rules for writing `skeleton` snippets

1. For each state variable, there is a corresponding component of the deterministic skeleton. The goal of such a snippet is to compute all the components.
2. When the skeleton is a map, the component corresponding to state variable `x` is named `Dx` and is the new value of `x` after one iteration of the map.
3. When the skeleton is a vectorfield, the component corresponding to state variable `x` is named `Dx` and is the value of dx/dt .
4. As with the other C snippets, all states, parameters and covariates, as well as the current time, `t`, will be defined in the context within which the snippet is executed.

Rules for writing `rprior` snippets

1. The goal of such a snippet is the replacement of parameters with values drawn from the prior distribution.
2. Within the context in which the snippet will be evaluated, only the parameters will be defined.
3. Hyperparameters can be included in the ordinary parameter list. Obviously, hyperparameters should not be replaced with random draws.

Rules for writing dprior snippets

1. The goal of such a snippet is computation of the prior probability density, or the log of same, at a given point in parameter space. This scalar value should be returned in the variable `lik`. When `give_log=1`, the user should return the log of the prior probability density.
2. Within the context in which the snippet will be evaluated, only the parameters and `give_log` will be defined.
3. Hyperparameters can be included in the ordinary parameter list. Obviously, hyperparameters should not be replaced with random draws.

Rules for writing parameter transformation snippets

1. The parameter transformation mapping a parameter vector from the scale used by the model codes to another scale is specified using the `toEstimationScale` argument whilst the transformation mapping a parameter vector from the alternative scale to that on which the model is defined is specified with the `fromEstimationScale` argument.
2. The goal of these snippets is the computation of the values of the transformed parameters. The value of transformed parameter `x` should be assigned to variable `Tx`.
3. Time-, state-, and covariate-dependent transformations are not allowed. Therefore, neither the time, nor any state variables, nor any of the covariates will be available in the context within which a parameter transformation snippet is executed.

Rules for writing initializer snippets

1. The goal of a this snippet is the construction of a state vector, i.e., the setting of the dynamical states at time t_0 .
2. In addition to the parameters and covariates (if any), the variable `t`, containing the time of the observation, will be defined in the context in which the snippet is executed.
3. **NB:** The `statenames` argument plays a particularly important role when the initializer is specified using a Csnippet. In particular, every state variable must be named in `statenames`. **Failure to follow this rule will result in undefined behavior.**

Viewing the generated C code

It can be useful to view the C code generated by calling `pomp` with one or more Csnippet arguments. To do so, set `options(verbose=TRUE)` before calling `pomp`. A message giving the name of the generated C file (in the session temporary directory) will be printed.

Author(s)

Aaron A. King

See Also

[pomp](#), [plugins](#), and the tutorials on the [package website](#).

dacca

Model of cholera transmission for historic Bengal.

Description

dacca is a pomp object containing census and cholera mortality data from the Dacca district of the former British province of Bengal over the years 1891 to 1940 together with a stochastic differential equation transmission model. The model is that of King et al. (2008). The parameters are the MLE for the SIRS model with seasonal reservoir.

Data are provided courtesy of Dr. Menno J. Bouma, London School of Tropical Medicine and Hygiene.

Details

dacca is a pomp object containing the model, data, and MLE parameters. Parameters that naturally range over the positive reals are log-transformed; parameters that range over the unit interval are logit-transformed; parameters that are naturally unbounded or take integer values are not transformed.

References

A. A. King, E. L. Ionides, M. Pascual, and M. J. Bouma, Inapparent infections and cholera dynamics, *Nature*, 454:877-880, 2008

See Also

[euler.sir](#), [pomp](#)

Examples

```
pompExample(dacca)
plot(dacca)
#MLEs on the natural scale
coef(dacca)
#MLEs on the transformed scale
coef(dacca,transform=TRUE)
plot(simulate(dacca))
# now change eps and simulate again
coef(dacca,"eps") <- 1
plot(simulate(dacca))
```

design

Design matrices for pomp calculations

Description

These functions are useful for generating designs for the exploration of parameter space. `sobolDesign` generate a Latin hypercube design using the Sobol' low-discrepancy sequence. `profileDesign` generates a data-frame where each row can be used as the starting point for a profile likelihood calculation. `sliceDesign` generates points along slices through a specified point.

Usage

```
sobolDesign(lower, upper, nseq)
profileDesign(..., lower, upper, nprof,
              stringsAsFactors = default.stringsAsFactors())
sliceDesign(center, ...)
```

Arguments

<code>lower, upper</code>	named numeric vectors giving the lower and upper bounds of the ranges, respectively.
<code>...</code>	In <code>profileDesign</code> , additional arguments specify the parameters over which to profile and the values of these parameters. In <code>sliceDesign</code> , additional numeric vector arguments specify the locations of points along the slices.
<code>nseq</code>	Total number of points requested.
<code>nprof</code>	The number of points per profile point.
<code>stringsAsFactors</code>	should character vectors be converted to factors?
<code>center</code>	<code>center</code> is a named numeric vector specifying the point through which the slice(s) is (are) to be taken.

Value

`sobolDesign`

`profileDesign` returns a data frame with `nprof` points per profile point. The other parameters in `vars` are sampled using `sobol`.

Author(s)

Aaron A. King

References

W. H. Press, S. A. Teukolsky, W. T. Vetterling, and B. P. Flannery, Numerical Recipes in C, Cambridge University Press, 1992

Examples

```
## Sobol low-discrepancy design
plot(sobolDesign(lower=c(a=0,b=100),upper=c(b=200,a=1),100))

## A one-parameter profile design:
x <- profileDesign(p=1:10,lower=c(a=0,b=0),upper=c(a=1,b=5),nprof=20)
dim(x)
plot(x)

## A two-parameter profile design:
x <- profileDesign(p=1:10,q=3:5,lower=c(a=0,b=0),upper=c(b=5,a=1),nprof=20)
dim(x)
plot(x)

## A single 11-point slice through the point c(A=3,B=8,C=0) along the B direction.
x <- sliceDesign(center=c(A=3,B=8,C=0),B=seq(0,10,by=1))
dim(x)
plot(x)

## Two slices through the same point along the A and C directions.
x <- sliceDesign(c(A=3,B=8,C=0),A=seq(0,5,by=1),C=seq(0,5,length=11))
dim(x)
plot(x)
```

eulermultinom

The Euler-multinomial distributions and Gamma white-noise processes

Description

This page documents both the Euler-multinomial family of distributions and the package's simulator of Gamma white-noise processes.

Usage

```
reulermultinom(n = 1, size, rate, dt)
deulermultinom(x, size, rate, dt, log = FALSE)
rgammawn(n = 1, sigma, dt)
```

Arguments

n	integer; number of random variates to generate.
size	scalar integer; number of individuals at risk.
rate	numeric vector of hazard rates.
sigma	numeric scalar; intensity of the Gamma white noise process.
dt	numeric scalar; duration of Euler step.
x	matrix or vector containing number of individuals that have succumbed to each death process.
log	logical; if TRUE, return logarithm(s) of probabilities.

Details

If N individuals face constant hazards of death in k ways at rates r_1, r_2, \dots, r_k , then in an interval of duration Δt , the number of individuals remaining alive and dying in each way is multinomially distributed:

$$(N - \sum_{i=1}^k \Delta n_i, \Delta n_1, \dots, \Delta n_k) \sim \text{multinomial}(N; p_0, p_1, \dots, p_k),$$

where Δn_i is the number of individuals dying in way i over the interval, the probability of remaining alive is $p_0 = \exp(-\sum_i r_i \Delta t)$, and the probability of dying in way j is

$$p_j = \frac{r_j}{\sum_i r_i} (1 - \exp(-\sum_i r_i \Delta t)).$$

In this case, we say that

$$(\Delta n_1, \dots, \Delta n_k) \sim \text{eulermultinom}(N, r, \Delta t),$$

where $r = (r_1, \dots, r_k)$. Draw m random samples from this distribution by doing

```
dn <- reulermultinom(n=m, size=N, rate=r, dt=dt),
```

where r is the vector of rates. Evaluate the probability that $x = (x_1, \dots, x_k)$ are the numbers of individuals who have died in each of the k ways over the interval $\Delta t = dt$, by doing

```
deulermultinom(x=x, size=N, rate=r, dt=dt).
```

Breto & Ionides (2011) discuss how an infinitesimally overdispersed death process can be constructed by compounding a binomial process with a Gamma white noise process. The Euler approximation of the resulting process can be obtained as follows. Let the increments of the equidispersed process be given by

```
reulermultinom(size=N, rate=r, dt=dt).
```

In this expression, replace the rate r with $r \Delta W / \Delta t$, where $\Delta W \sim \text{Gamma}(\Delta t / \sigma^2, \sigma^2)$ is the increment of an integrated Gamma white noise process with intensity σ . That is, ΔW has mean Δt and variance $\sigma^2 \Delta t$. The resulting process is overdispersed and converges (as Δt goes to zero) to a well-defined process. The following lines of R code accomplish this:

```
dW <- rgammawn(sigma=sigma, dt=dt)
```

```
dn <- reulermultinom(size=N, rate=r, dt=dW)
```

or

```
dn <- reulermultinom(size=N, rate=r*dW/dt, dt=dt).
```

He et al. use such overdispersed death processes in modeling measles.

For all of the functions described here, access to the underlying C routines is available: see below.

Value

<code>reulermultinom</code>	Returns a <code>length(rate)</code> by <code>n</code> matrix. Each column is a different random draw. Each row contains the numbers of individuals succumbed to the corresponding process.
<code>deulermultinom</code>	Returns a vector (of length equal to the number of columns of <code>x</code>) containing the probabilities of observing each column of <code>x</code> given the specified parameters (<code>size</code> , <code>rate</code> , <code>dt</code>).
<code>rgammawn</code>	Returns a vector of length <code>n</code> containing random increments of the integrated Gamma white noise process with intensity <code>sigma</code> .

C API

An interface for C codes using these functions is provided by the package. At an R prompt, execute

```
file.show(system.file("include/pomp.h", package="pomp"))
```

to view the header file that defines and explains the API.

Author(s)

Aaron A. King

References

- C. Breto & E. L. Ionides, Compound Markov counting processes and their applications to modeling infinitesimally over-dispersed systems. *Stoch. Proc. Appl.*, 121:2571–2591, 2011.
- D. He, E. L. Ionides, & A. A. King, Plug-and-play inference for disease dynamics: measles in large and small populations as a case study. *J. R. Soc. Interface*, 7:271–283, 2010.

Examples

```
print(dn <- reulermultinom(5,size=100,rate=c(a=1,b=2,c=3),dt=0.1))
deulermultinom(x=dn,size=100,rate=c(1,2,3),dt=0.1)
## an Euler-multinomial with overdispersed transitions:
dt <- 0.1
dW <- rgammawn(sigma=0.1,dt=dt)
print(dn <- reulermultinom(5,size=100,rate=c(a=1,b=2,c=3),dt=dW))
```

Example pomp models *Examples of the construction of POMP models*

Description

`pompExample` loads pre-built example pomp objects.

Usage

```
pompExample(example, ..., envir = .GlobalEnv)
```

Arguments

example	example to load given as a name or literal character string. Evoked without an argument, <code>pompExample</code> lists all available examples.
...	additional arguments define symbols in the environment within which the example code is executed.
envir	the environment into which the objects should be loaded. If <code>envir=NULL</code> , then the created objects are returned in a list.

Details

Directories listed in the global option `pomp.examples` (which can be changed using `options()`) are searched for file named '`<example>.R`'. If found, this file will be sourced in a temporary environment. Additional arguments to `pompExample` define variables within this environment and will therefore be available when the code in '`<example>.R`' is sourced.

The codes that construct these pomp objects can be found in the 'examples' directory in the installed package. Do `system.file("examples", package="pomp")` to find this directory.

Value

By default, `pompExample` has the side effect of creating one or more objects in the global workspace. If `envir=NULL`, there are no side effects; rather, the objects are returned as a list.

Author(s)

Aaron A. King

See Also

[blowflies](#), [dacca](#), [gompertz](#), [ou2](#), [ricker](#), [rw2](#), [euler.sir](#), [gillespie.sir](#), [bbs](#)

Examples

```
pompExample()
pompExample(euler.sir)
pompExample("gompertz")
pompExample(ricker,envir=NULL)
## Not run:
  file.show(system.file("examples/bbs.R", package="pomp"))

## End(Not run)
```

gompertz

*Gompertz model with log-normal observations.***Description**

gompertz is a pomp object encoding a stochastic Gompertz population model with log-normal measurement error.

Details

The state process is $X_{t+1} = K^{1-S} X_t^S \epsilon_t$, where $S = e^{-r}$ and the ϵ_t are i.i.d. lognormal random deviates with variance σ^2 . The observed variables Y_t are distributed as $\text{lognormal}(\log X_t, \tau)$. Parameters include the per-capita growth rate r , the carrying capacity K , the process noise s.d. σ , the measurement error s.d. τ , and the initial condition X_0 . The pomp object includes parameter transformations that log-transform the parameters for estimation purposes.

See Also

pomp, ricker, and the tutorials at <http://kingaa.github.io/pomp>.

Examples

```
pompExample(gompertz)
plot(gompertz)
coef(gompertz)
coef(gompertz, transform=TRUE)
```

Iterated filtering

*Maximum likelihood by iterated filtering***Description**

Iterated filtering algorithms for estimating the parameters of a partially-observed Markov process. Running `mif` causes the iterated filtering algorithm to run for a specified number of iterations. At each iteration, the particle filter is performed on a perturbed version of the model. Specifically, parameters to be estimated are subjected to random perturbations at each observation. This extra variability effectively smooths the likelihood surface and combats particle depletion by introducing diversity into the population of particles. As the iterations progress, the magnitude of the perturbations is diminished according to a user-specified cooling schedule. For most purposes, `mif` has been superseded by `mif2`.

Usage

```
## S4 method for signature pomp
mif(object, Nmif = 1, start, ivps = character(0),
    particles, rw.sd, Np, ic.lag, var.factor = 1,
    cooling.type, cooling.fraction.50,
    method = c("mif", "unweighted", "fp", "mif2"),
    tol = 1e-17, max.fail = Inf,
    verbose = getOption("verbose"), transform = FALSE, ...)
## S4 method for signature pfilterd.pomp
mif(object, Nmif = 1, Np, tol, ...)
## S4 method for signature mif
mif(object, Nmif, start, ivps,
    particles, rw.sd, Np, ic.lag, var.factor,
    cooling.type, cooling.fraction.50,
    method, tol, transform, ...)
## S4 method for signature mif
continue(object, Nmif = 1, ...)
## S4 method for signature mif
conv.rec(object, pars, transform = FALSE, ...)
## S4 method for signature mifList
conv.rec(object, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>Nmif</code>	The number of filtering iterations to perform.
<code>start</code>	named numerical vector; the starting guess of the parameters.
<code>ivps</code>	optional character vector naming the initial-value parameters (IVPs) to be estimated. Every parameter named in <code>ivps</code> must have a positive random-walk standard deviation specified in <code>rw.sd</code> . If there are no regular parameters with positive <code>rw.sd</code> , i.e., only IVPs are to be estimated, see below “Using <code>mif</code> to estimate initial-value parameters only”.
<code>particles</code>	Function of prototype <code>particles(Np, center, sd, ...)</code> which sets up the starting particle matrix by drawing a sample of size <code>Np</code> from the starting particle distribution centered at <code>center</code> and of width <code>sd</code> . If <code>particles</code> is not supplied by the user, the default behavior is to draw the particles from a multivariate normal distribution with mean <code>center</code> and standard deviation <code>sd</code> .
<code>rw.sd</code>	numeric vector with names; the intensity of the random walk to be applied to parameters. The random walk is only applied to parameters named in <code>pars</code> (i.e., not to those named in <code>ivps</code>). The algorithm requires that the random walk be nontrivial, so each element in <code>rw.sd[pars]</code> must be positive. <code>rw.sd</code> is also used to scale the initial-value parameters (via the <code>particles</code> function). Therefore, each element of <code>rw.sd[ivps]</code> must be positive. The following must be satisfied: <code>names(rw.sd)</code> must be a subset of <code>names(start)</code> , <code>rw.sd</code> must be non-negative (zeros are simply ignored), the name of every positive element of <code>rw.sd</code> must be in either <code>pars</code> or <code>ivps</code> .

<code>Np</code>	the number of particles to use in filtering. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timestep, one may specify <code>Np</code> either as a vector of positive integers (of length <code>length(time(object, t0=TRUE))</code>) or as a function taking a positive integer argument. In the latter case, <code>Np(k)</code> must be a single positive integer, representing the number of particles to be used at the k -th timestep: <code>Np(0)</code> is the number of particles to use going from <code>timezero(object)</code> to <code>time(object)[1]</code> , <code>Np(1)</code> , from <code>timezero(object)</code> to <code>time(object)[1]</code> , and so on, while when <code>T=length(time(object, t0=TRUE))</code> , <code>Np(T)</code> is the number of particles to sample at the end of the time-series.
<code>ic.lag</code>	a positive integer; the timepoint for fixed-lag smoothing of initial-value parameters. The mif update for initial-value parameters consists of replacing them by their filtering mean at time <code>times[ic.lag]</code> , where <code>times=time(object)</code> . It makes no sense to set <code>ic.lag>length(times)</code> ; if it is so set, <code>ic.lag</code> is set to <code>length(times)</code> with a warning. For <code>method="mif2"</code> , the default is <code>ic.lag=length(times)</code> .
<code>var.factor</code>	optional positive scalar; the scaling coefficient relating the width of the starting particle distribution to <code>rw.sd</code> . In particular, the width of the distribution of particles at the start of the first mif iteration will be <code>random.walk.sd*var.factor</code> . By default, <code>var.factor=1</code> .
<code>cooling.type</code> , <code>cooling.fraction.50</code>	specifications for the cooling schedule, i.e., the manner in which the intensity of the parameter perturbations is reduced with successive filtering iterations. <code>cooling.type</code> specifies the nature of the cooling schedule. When <code>cooling.type="geometric"</code> , on the n -th mif iteration, the relative perturbation intensity is <code>cooling.fraction.50^(n/50)</code> . When <code>cooling.type="hyperbolic"</code> , on the n -th mif iteration, the relative perturbation intensity is $(s+1)/(s+n)$, where $(s+1)/(s+50)=\text{cooling.fraction.50}$. <code>cooling.fraction.50</code> is the relative magnitude of the parameter perturbations after 50 mif iterations.
<code>method</code>	<code>method</code> sets the update rule used in the algorithm. <code>method="mif"</code> uses the iterated filtering update rule (Ionides 2006, 2011); <code>method="unweighted"</code> updates the parameter to the unweighted average of the filtering means of the parameters at each time; <code>method="fp"</code> updates the parameter to the filtering mean at the end of the time series. <code>method="mif2"</code> implements an incomplete version of the iterated Bayes map method of Ionides (2015). The latter method is, by every indication, both more efficient and more stable.
<code>tol</code> , <code>max.fail</code>	See the description under pfilter .
<code>verbose</code>	logical; if TRUE, print progress reports.
<code>transform</code>	logical; if TRUE, optimization is performed on the transformed scale, as defined by the user-supplied parameter transformations (see pomp).
<code>...</code>	additional arguments that override the defaults.
<code>pars</code>	names of parameters.

Value

Upon successful completion, `mif` returns an object of class `mif`. The latter inherits from the `pfilterd.pomp` and `pomp` classes.

IF2

A more full-featured version of the improved iterated filtering algorithm (IF2) is implemented as `mif2`.

Regular parameters vs initial-value parameters

Initial-value parameters (IVPs) differ from regular parameters in that the majority of the information about these parameters is restricted to the early part of the time series. That is, increasing the length of the time series provides progressively less additional information about IVPs than it does about regular parameters. In `mif`, while regular parameters are perturbed at the initial time and after every observation, IVPs are perturbed only at the initial time.

Re-running `mif` Iterations

To re-run a sequence of `mif` iterations, one can use the `mif` method on a `mif` object. By default, the same parameters used for the original `mif` run are re-used (except for `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing `mif` Iterations

One can resume a series of `mif` iterations from where one left off using the `continue` method. A call to `mif` to perform `Nmif=m` iterations followed by a call to `continue` to perform `Nmif=n` iterations will produce precisely the same effect as a single call to `mif` to perform `Nmif=m+n` iterations. By default, all the algorithmic parameters are the same as used in the original call to `mif`. Additional arguments will override the defaults.

Using `mif` to estimate initial-value parameters only

One can use `mif`'s fixed-lag smoothing to estimate only initial value parameters (IVPs). In this case, `pars` is left empty and the IVPs to be estimated are named in `ivps`. If `theta` is the current parameter vector, then at each `mif` iteration, `Np` particles are drawn from a distribution centered at `theta` and with width proportional to `var.factor*rw.sd`, a particle filtering operation is performed, and `theta` is replaced by the filtering mean at `time(object)[ic.lag]`. Note the implication that, when `mif` is used in this way on a time series any longer than `ic.lag`, unnecessary work is done. If the time series in `object` is longer than `ic.lag`, consider replacing `object` with `window(object,end=ic.lag)`.

Methods

Methods that can be used to manipulate, display, or extract information from a `mif` object:

conv.rec `conv.rec(object, pars = NULL)` returns the columns of the convergence-record matrix corresponding to the names in `pars`. By default, all rows are returned.

logLik Returns the value in the loglik slot. NB: this is *not* the same as the likelihood of the model at the MLE!

c Concatenates mif objects into a mifList.

plot Plots a series of diagnostic plots when applied to a mif or mifList object.

Details

If particles is not specified, the default behavior is to draw the particles from a multivariate normal distribution. **It is the user's responsibility to ensure that, if the optional particles argument is given, that the particles function satisfies the following conditions:**

particles has at least the following arguments: Np, center, sd, and . . . Np may be assumed to be a positive integer; center and sd will be named vectors of the same length. Additional arguments may be specified; these will be filled with the elements of the userdata slot of the underlying pomp object (see [pomp](#)).

particles returns a length(center) x Np matrix with rownames matching the names of center and sd. Each column represents a distinct particle.

The center of the particle distribution returned by particles should be center. The width of the particle distribution should vary monotonically with sd. In particular, when sd=0, the particles should return matrices with Np identical columns, each given by the parameters specified in center.

Author(s)

Aaron A. King

References

- E. L. Ionides, C. Breto, & A. A. King, Inference for nonlinear dynamical systems, Proc. Natl. Acad. Sci. U.S.A., 103:18438–18443, 2006.
- E. L. Ionides, A. Bhadra, Y. Atchad'e, & A. A. King, Iterated filtering, Annals of Statistics, 39:1776–1802, 2011.
- E. L. Ionides, D. Nguyen, Y. Atchad'e, S. Stoev, and A. A. King. Inference for dynamic and latent variable models via iterated, perturbed Bayes maps. Proc. Natl. Acad. Sci. U.S.A., 112:719–724, 2015.
- A. A. King, E. L. Ionides, M. Pascual, and M. J. Bouma, Inapparent infections and cholera dynamics, Nature, 454:877–880, 2008.

See Also

[pomp](#), [pfilter](#), [mif2](#)

Description

An improved iterated filtering algorithm for estimating the parameters of a partially-observed Markov process. Running `mif2` causes the algorithm to perform a specified number of particle-filter iterations. At each iteration, the particle filter is performed on a perturbed version of the model, in which the parameters to be estimated are subjected to random perturbations at each observation. This extra variability effectively smooths the likelihood surface and combats particle depletion by introducing diversity into particle population. As the iterations progress, the magnitude of the perturbations is diminished according to a user-specified cooling schedule. The algorithm is presented and justified in Ionides et al. (2015).

Usage

```
## S4 method for signature pomp
mif2(object, Nmif = 1, start, Np, rw.sd, transform = FALSE,
      cooling.type = c("hyperbolic", "geometric"), cooling.fraction.50,
      tol = 1e-17, max.fail = Inf, verbose = getOption("verbose"), ...)
## S4 method for signature pfilterd.pomp
mif2(object, Nmif = 1, Np, tol, ...)
## S4 method for signature mif2d.pomp
mif2(object, Nmif, start, Np, rw.sd, transform,
      cooling.type, cooling.fraction.50, tol, ...)
## S4 method for signature mif2d.pomp
continue(object, Nmif = 1, ...)
## S4 method for signature mif2d.pomp
conv.rec(object, pars, transform = FALSE, ...)
## S4 method for signature mif2List
conv.rec(object, ...)
rw.sd(...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>Nmif</code>	The number of filtering iterations to perform.
<code>start</code>	named numerical vector; the starting guess of the parameters. By default, <code>start=coef(object)</code> .
<code>Np</code>	the number of particles to use in filtering. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timestep, one may specify <code>Np</code> either as a vector of positive integers (of length <code>length(time(object))</code>) or as a function taking a positive integer argument. In

the latter case, `Np(n)` must be a single positive integer, representing the number of particles to be used at the n -th timestep: `Np(1)` is the number of particles to use going from `timezero(object)` to `time(object)[1]`, `Np(2)`, from `time(object)[1]` to `time(object)[2]`, and so on. **Note that this behavior differs from that of `mif`!**

<code>rw.sd</code>	<p>specification of the magnitude of the random-walk perturbations that will be applied to some or all model parameters. Parameters that are to be estimated should have positive perturbations specified here. The specification is given using the <code>rw.sd</code> function, which creates a list of unevaluated expressions. The latter are evaluated in a context where the model time variable is defined (as <code>time</code>). The expression <code>ivp(s)</code> can be used in this context as shorthand for <code>ifelse(time==time[1],s,0)</code>.</p> <p>Likewise, <code>ivp(s,lag)</code> is equivalent to <code>ifelse(time==time[lag],s,0)</code>.</p> <p>See below for some examples. The perturbations that are applied are normally distributed with the specified s.d. If <code>transform = TRUE</code>, then they are applied on the estimation scale.</p>
<code>transform</code>	logical; if <code>TRUE</code> , optimization is performed on the estimation scale, as defined by the user-supplied parameter transformations (see pomp). This can be used, for example, to enforce positivity or interval constraints on model parameters. See the tutorials on the package website for examples.
<code>cooling.type</code> , <code>cooling.fraction.50</code>	<p>specifications for the cooling schedule, i.e., the manner in which the intensity of the parameter perturbations is reduced with successive filtering iterations. <code>cooling.type</code> specifies the nature of the cooling schedule. See below (under “Specifying the perturbations”) for more detail.</p>
<code>tol</code> , <code>max.fail</code>	passed to the particle filter. See the descriptions under pfilter .
<code>verbose</code>	logical; if <code>TRUE</code> , print progress reports.
<code>...</code>	additional arguments that override the defaults.
<code>pars</code>	names of parameters.

Value

Upon successful completion, `mif2` returns an object of class `mif2d.pomp`. This class inherits from the [pfilterd.pomp](#) and [pomp](#) classes.

Specifying the perturbations: the `rw.sd` function

This function simply returns a list containing its arguments as unevaluated expressions. These are then evaluated in a context containing the model time variable. This allows for easy specification of the structure of the perturbations that are to be applied. For example,

```
rw.sd(a=0.05,
      b=ifelse(0.2,time==time[1],0),
      c=ivp(0.2),
```

```
d=ifelse(time==time[13],0.2,0),
e=ivp(0.2,lag=13),
f=ifelse(time<23,0.02,0))
```

results in perturbations of parameter a with s.d. 0.05 at every time step, while parameters b and c both get perturbations of s.d. 0.2 only before the first observation. Parameters d and e , by contrast, get perturbations of s.d. 0.2 only before the thirteenth observation. Finally, parameter f gets a random perturbation of size 0.02 before every observation falling before $t = 23$.

On the m -th IF2 iteration, prior to time-point n , the d -th parameter is given a random increment normally distributed with mean 0 and standard deviation $c_{m,n}\sigma_{d,n}$, where c is the cooling schedule and σ is specified using `rw.sd`, as described above. Let N be the length of the time series and $\alpha = \text{cooling.fraction}/50$. Then, when `cooling.type="geometric"`, we have

$$c_{m,n} = \alpha^{\frac{n-1+(m-1)N}{50N}}.$$

When `cooling.type="hyperbolic"`, we have

$$c_{m,n} = \frac{s+1}{s+n+(m-1)N},$$

where s satisfies

$$\frac{s+1}{s+50N} = \alpha.$$

Thus, in either case, the perturbations at the end of 50 IF2 iterations are a fraction α smaller than they are at first.

Re-running mif2 Iterations

To re-run a sequence of `mif2` iterations, one can use the `mif2` method on a `mif2d.pomp` object. By default, the same parameters used for the original `mif2` run are re-used (except for `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing mif2 Iterations

One can resume a series of `mif2` iterations from where one left off using the `continue` method. A call to `mif2` to perform $N_{\text{mif}}=m$ iterations followed by a call to `continue` to perform $N_{\text{mif}}=n$ iterations will produce precisely the same effect as a single call to `mif2` to perform $N_{\text{mif}}=m+n$ iterations. By default, all the algorithmic parameters are the same as used in the original call to `mif2`. Additional arguments will override these defaults.

Methods

Methods that can be used to manipulate, display, or extract information from a `mif2d.pomp` object:

conv.rec `conv.rec(object, pars = NULL)` returns the columns of the convergence-record matrix corresponding to the names in `pars`. By default, all rows are returned.

logLik Returns the value in the `loglik` slot. NB: this is *not* the same as the likelihood of the model at the MLE!

c Concatenates `mif2d.pomp` objects into a `mif2List`.

plot Plots a series of diagnostic plots when applied to a `mif2d.pomp` or `mif2List` object.

Author(s)

Aaron A. King, Edward L. Ionides, and Dao Nguyen

References

E. L. Ionides, D. Nguyen, Y. Atchad'e, S. Stoev, and A. A. King. Inference for dynamic and latent variable models via iterated, perturbed Bayes maps. Proc. Natl. Acad. Sci. U.S.A., 112:719–724, 2015.

See Also

[pomp](#), [pfilter](#), [mif](#), and the tutorials on the [package website](#).

Examples

```
## Not run:
pompExample(ou2)

guess1 <- guess2 <- coef(ou2)
guess1[c(x1.0,x2.0,alpha.2,alpha.3)] <- 0.5*guess1[c(x1.0,x2.0,alpha.2,alpha.3)]
guess2[c(x1.0,x2.0,alpha.2,alpha.3)] <- 1.5*guess1[c(x1.0,x2.0,alpha.2,alpha.3)]

m1 <- mif2(ou2,Nmif=100,start=guess1,Np=1000,
           cooling.type="hyperbolic",cooling.fraction.50=0.05,
           rw.sd=rw.sd(x1.0=ivp(0.5),x2.0=ivp(0.5),
                       alpha.2=0.1,alpha.3=0.1))

m2 <- mif2(ou2,Nmif=100,start=guess2,Np=1000,
           cooling.type="hyperbolic",cooling.fraction.50=0.05,
           rw.sd=rw.sd(x1.0=ivp(0.5),x2.0=ivp(0.5),
                       alpha.2=0.1,alpha.3=0.1))

plot(c(m1,m2))

rbind(mle1=c(coef(m1),loglik=logLik(pfilter(m1,Np=1000))),
      mle2=c(coef(m2),loglik=logLik(pfilter(m1,Np=1000))),
      truth=c(coef(ou2),loglik=logLik(pfilter(m1,Np=1000))))

## End(Not run)
```

logmeanexp

The log-mean-exp trick

Description

logmeanexp computes

$$\log \frac{1}{N} \sum_{n=1}^N e_i^x,$$

avoiding over- and under-flow in doing so. It can optionally return an estimate of the standard error in this quantity.

Usage

```
logmeanexp(x, se = FALSE)
```

Arguments

x	numeric
se	logical; give approximate standard error?

Details

When `se = TRUE`, `logmeanexp` uses a jackknife estimate of the variance in $\log(x)$.

Value

$\log(\text{mean}(\exp(x)))$ computed so as to avoid over- or underflow. If `se = FALSE`, the approximate standard error is returned as well.

Author(s)

Aaron A. King

Examples

```
## generate a bifurcation diagram for the Ricker map
pompExample(ricker)
ll <- replicate(n=5, logLik(pfilter(ricker, Np=1000)))
## an estimate of the log likelihood:
logmeanexp(ll)
## with standard error:
logmeanexp(ll, se=TRUE)
```

Low-level-interface *pomp low-level interface*

Description

A `pomp` object implements a partially observed Markov process (POMP) model. Basic operations on this model (with shorthand terms) include:

1. simulation of the state process given parameters (`rprocess`)
2. evaluation of the likelihood of a given state trajectory given parameters (`dprocess`)
3. simulation of the observation process given the states and parameters (`rmeasure`)
4. evaluation of the likelihood of a set of observations given the states and parameters (`dmeasure`)

5. simulation from the prior probability distribution (rprior)
6. evaluation of the prior probability density (dprior)
7. simulation from the distribution of initial states, given parameters (init.state)
8. evaluation of the deterministic skeleton at a point in state space, given parameters (skeleton)
9. computation of a trajectory of the deterministic skeleton given parameters (trajectory)

pomp provides S4 methods that implement each of these basic operations. These operations can be combined to implement computations and statistical inference methods that depend only on a model's POMP structure. For convenience, parameter transformations may also be enclosed in a pomp object.

This page documents these elements.

Usage

```
## S4 method for signature pomp
rprocess(object, xstart, times, params, offset = 0, ...)
## S4 method for signature pomp
dprocess(object, x, times, params, log = FALSE, ...)
## S4 method for signature pomp
rmeasure(object, x, times, params, ...)
## S4 method for signature pomp
dmeasure(object, y, x, times, params, log = FALSE, ...)
## S4 method for signature pomp
dprior(object, params, log = FALSE, ...)
## S4 method for signature pomp
rprior(object, params, ...)
## S4 method for signature pomp
init.state(object, params, t0, ...)
## S4 method for signature pomp
skeleton(object, x, t, params, ...)
## S4 method for signature pomp
trajectory(object, params, times, t0, as.data.frame = FALSE, ...)
## S4 method for signature pomp
pompLoad(object, ...)
## S4 method for signature pomp
pompUnload(object, ...)
```

Arguments

object	an object of class pomp.
xstart	an nvar x nrep matrix containing the starting state of the system. Columns of xstart correspond to states; rows to components of the state vector. One independent simulation will be performed for each column. Note that in this case, params must also have nrep columns.
x	a rank-3 array containing states of the unobserved process. The dimensions of x are nvars x nrep x ntimes, where nvars is the number of state variables, nrep is the number of replicates, and ntimes is the length of times.

<code>y</code>	a matrix containing observations. The dimensions of <code>y</code> are <code>nobs</code> x <code>ntimes</code> , where <code>nobs</code> is the number of observables and <code>ntimes</code> is the length of times.
<code>times, t</code>	a numeric vector (length <code>ntimes</code>) containing times. These must be in non-decreasing order.
<code>params</code>	a <code>npar</code> x <code>nrep</code> matrix of parameters. Each column is an independent parameter set and is paired with the corresponding column of <code>x</code> or <code>xstart</code> . In the case of <code>init.state</code> , <code>params</code> is a named vector of parameters.
<code>offset</code>	integer; the first <code>offset</code> times in <code>times</code> will not be returned.
<code>t0</code>	the initial time at which initial states are requested.
<code>log</code>	if TRUE, log probabilities are returned.
<code>as.data.frame</code>	logical; if TRUE, return the result as a data-frame.
<code>...</code>	In trajectory, additional arguments are passed to the ODE integrator (if the skeleton is a vectorfield) and ignored if it is a map. See ode for a description of the additional arguments accepted. In all other cases, additional arguments are ignored.

rprocess

`rprocess` simulates the process-model portion of partially-observed Markov process.

When `rprocess` is called, the first entry of `times` is taken to be the initial time (i.e., that corresponding to `xstart`). Subsequent times are the additional times at which the state of the simulated processes are required.

`rprocess` returns a rank-3 array with rownames. Suppose `x` is the array returned. Then

```
dim(x)=c(nvars,nrep,ntimes-offset),
```

where `nvars` is the number of state variables (`=nrow(xstart)`), `nrep` is the number of independent realizations simulated (`=ncol(xstart)`), and `ntimes` is the length of the vector `times`. `x[,j,k]` is the value of the state process in the `j`-th realization at time `times[k+offset]`. The rownames of `x` must correspond to those of `xstart`.

dprocess

`dprocess` evaluates the probability density of a sequence of consecutive state transitions.

`dprocess` returns a matrix of dimensions `nrep` x `ntimes-1`. If `d` is the returned matrix, `d[j,k]` is the likelihood of the transition from state `x[,j,k-1]` at time `times[k-1]` to state `x[,j,k]` at time `times[k]`.

rmeasure

`rmeasure` simulate the measurement model given states and parameters.

`rmeasure` returns a rank-3 array of dimensions `nobs` x `nrep` x `ntimes`, where `nobs` is the number of observed variables.

dmeasure

dmeasure evaluates the probability density of observations given states.

dmeasure returns a matrix of dimensions nreps x ntimes. If d is the returned matrix, $d[j,k]$ is the likelihood of the observation $y[j,k]$ at time $times[k]$ given the state $x[j,k]$.

dprior, rprior

dprior evaluates the prior probability density and rprior simulates from the prior.

init.state

init.state returns an nvar x nrep matrix of state-process initial conditions when given an nvar x nrep matrix of parameters, params, and an initial time t0. By default, t0 is the initial time defined when the pomp object was constructed.

skeleton

The method skeleton evaluates the deterministic skeleton at a point or points in state space, given parameters. In the case of a discrete-time system, the skeleton is a map. In the case of a continuous-time system, the skeleton is a vectorfield. NB: skeleton just evaluates the deterministic skeleton; it does not iterate or integrate.

skeleton returns an array of dimensions nvar x nrep x ntimes. If f is the returned matrix, $f[i,j,k]$ is the i-th component of the deterministic skeleton at time $times[k]$ given the state $x[j,k]$ and parameters $params[j]$.

trajectory

trajectory computes a trajectory of the deterministic skeleton of a Markov process. In the case of a discrete-time system, the deterministic skeleton is a map and a trajectory is obtained by iterating the map. In the case of a continuous-time system, the deterministic skeleton is a vectorfield; trajectory uses the numerical solvers in [deSolve](#) to integrate the vectorfield.

trajectory returns an array of dimensions nvar x nrep x ntimes. If x is the returned matrix, $x[i,j,k]$ is the i-th component of the state vector at time $times[k]$ given parameters $params[j]$.

When the skeleton is a vectorfield, trajectory integrates it using [ode](#). When the skeleton is a map, trajectory iterates it. By default, time is advanced 1 unit per iteration. The user can change this behavior by specifying the desired timestep using the argument `skelmap.delta.t` in the construction of the pomp object.

Parameter transformations

User-defined parameter transformations enclosed in the pomp object can be accessed via [partrans](#).

pompLoad, pompUnload

pompLoad and pompUnload cause compiled codes associated with object to be dynamically linked or unlinked, respectively. When Csnippets are used in the construction of a pomp object, the resulting shared-object library is dynamically loaded (linked) before each use, and unloaded afterward.

These functions are provided because in some instances, greater control may be desired. These functions have no effect on shared-object libraries linked by the user.

Author(s)

Aaron A. King

See Also

[pomp](#), [pomp methods](#)

Examples

```
pompExample(ricker)

p <- parmat(c(r=42,phi=10,sigma=0.3,N.0=7,e.0=0),10)
t <- c(1:10,20,30)
t0 <- 0
x0 <- init.state(ricker,params=p,t0=t0)
x <- rprocess(ricker,xstart=x0,times=c(t0,t),params=p,offset=1)
y <- rmeasure(ricker,params=p,x=x,times=t)
ll <- dmeasure(ricker,y=y[,3,,drop=FALSE],x=x,times=t,params=p,log=TRUE)
apply(ll,1,sum)
f <- skeleton(ricker,x=x,t=t,params=p)
z <- trajectory(ricker,params=p,times=t,t0=t0)

## short arguments are recycled:
p <- c(r=42,phi=10,sigma=0.3,N.0=7,e.0=0)
t <- c(1:10,20,30)
t0 <- 0
x0 <- init.state(ricker,params=p,t0=t0)
x <- rprocess(ricker,xstart=x0,times=c(t0,t),params=p,offset=1)
y <- rmeasure(ricker,params=p,x=x,times=t)
ll <- dmeasure(ricker,y=y,x=x,times=t,params=p,log=TRUE)
f <- skeleton(ricker,x=x,t=t,params=p)
z <- trajectory(ricker,params=p,times=t,t0=t0)
```

MCMC proposal distributions

MCMC proposal distributions

Description

Functions to construct proposal distributions for use with MCMC methods.

Usage

```
mvn.rw(rw.var)
mvn.diag.rw(rw.sd)
mvn.rw.adaptive(rw.sd, rw.var, scale.start = NA, scale.cooling = 0.999,
               shape.start = NA, target = 0.234, max.scaling = 50)
```

Arguments

`rw.var` square numeric matrix with row- and column-names. Specifies the variance-covariance matrix for a multivariate normal random-walk proposal distribution.

`rw.sd` named numeric vector; random-walk SDs for a multivariate normal random-walk proposal with diagonal variance-covariance matrix.

`scale.start`, `scale.cooling`, `shape.start`, `target`, `max.scaling` parameters to control the proposal adaptation algorithm. Beginning with MCMC iteration `scale.start`, the scale of the proposal covariance matrix will be adjusted in an effort to match the target acceptance ratio. The parameters `scale.cooling` and `max.scaling` adjust the scale adaptation. Beginning with iteration `shape.start`, a scaled empirical covariance matrix will be used for the proposals.

Value

Each of these constructs a function suitable for use as the proposal argument of `pmmcmc` or `abc`. Given a parameter vector, each such function returns a single draw from the corresponding proposal distribution.

Author(s)

Aaron A. King

See Also

[pmmcmc](#), [abc](#)

Nonlinear forecasting *Parameter estimation my maximum simulated quasi-likelihood (non-linear forecasting)*

Description

`nlf` calls an optimizer to maximize the nonlinear forecasting (NLF) goodness of fit. The latter is computed by simulating data from a model, fitting a nonlinear autoregressive model to the simulated time series, and quantifying the ability of the resulting fitted model to predict the data time series. NLF is an ‘indirect inference’ method using a quasi-likelihood as the objective function.

Usage

```
## S4 method for signature pomp
nlf(object, start, est, lags, period = NA, tensor = FALSE,
     nconverge=1000, nasymp=1000, seed = 1066,
     transform.data, nrbf = 4,
     method = c("subplex", "Nelder-Mead", "BFGS", "CG", "L-BFGS-B", "SANN", "Brent"),
     skip.se = FALSE, verbose = getOption("verbose"),
     bootstrap=FALSE, bootsamp = NULL,
```

```

    lql.frac = 0.1, se.par.frac = 0.1, eval.only = FALSE,
    transform = FALSE, ...)
## S4 method for signature nlfd.pomp
nlf(object, start, est, lags, period, tensor,
    nconverge, nasymp, seed, transform.data, nrbf, method,
    lql.frac, se.par.frac, transform, ...)

```

Arguments

<code>object</code>	A pomp object, with the data and model to fit to it.
<code>start</code>	Named numeric vector with guessed parameters.
<code>est</code>	Vector containing the names or indices of parameters to be estimated.
<code>lags</code>	A vector specifying the lags to use when constructing the nonlinear autoregressive prediction model. The first lag is the prediction interval.
<code>period</code>	numeric; <code>period=NA</code> means the model is nonseasonal. <code>period>0</code> is the period of seasonal forcing in 'real time'.
<code>tensor</code>	logical; if <code>FALSE</code> , the fitted model is a generalized additive model with time mod period as one of the predictors, i.e., a gam with time-varying intercept. If <code>TRUE</code> , the fitted model is a gam with lagged state variables as predictors and time-periodic coefficients, constructed using tensor products of basis functions of state variables with basis functions of time.
<code>nconverge</code>	Number of convergence timesteps to be discarded from the model simulation.
<code>nasymp</code>	Number of asymptotic timesteps to be recorded from the model simulation.
<code>seed</code>	Integer specifying the random number seed to use. When fitting, it is usually best to always run the simulations with the same sequence of random numbers, which is accomplished by setting <code>seed</code> to an integer. If you want a truly random simulation, set <code>seed=NULL</code> .
<code>transform</code>	logical; if <code>TRUE</code> , parameters are optimized on the transformed scale.
<code>transform.data</code>	optional function. If specified, forecasting is performed using data and model simulations transformed by this function. By default, <code>transform.data</code> is the identity function, i.e., no transformation is performed. The main purpose of <code>transform.data</code> is to achieve approximately multivariate normal forecasting errors. If data are univariate, <code>transform.data</code> should take a scalar and return a scalar. If data are multivariate, <code>transform.data</code> should assume a vector input and return a vector of the same length.
<code>nrbf</code>	A scalar specifying the number of radial basis functions to be used at each lag.
<code>method</code>	Optimization method. Choices are subplex and any of the methods used by optim .
<code>skip.se</code>	Logical; if <code>TRUE</code> , skip the computation of standard errors.
<code>verbose</code>	Logical; if <code>TRUE</code> , the negative log quasiliikelihood and parameter values are printed at each iteration of the optimizer.
<code>bootstrap</code>	Logical; if <code>TRUE</code> the indices in <code>bootsamp</code> will determine which of the conditional likelihood values be used in computing the quasi-loglikelihood.

<code>bootsamp</code>	Vector of integers; used to have the quasi-loglikelihood evaluated using a bootstrap re-sampling of the data set.
<code>lql.frac</code>	target fractional change in log quasi-likelihood for quadratic standard error estimate
<code>se.par.frac</code>	initial parameter-change fraction for quadratic standard error estimate
<code>eval.only</code>	logical; if TRUE, no optimization is attempted and the quasi-loglikelihood value is evaluated at the start parameters.
<code>...</code>	Arguments that will be passed to <code>optim</code> or <code>subplex</code> in the control list.

Details

This runs an optimizer to maximize `nlf.objfun`.

Value

An object of class `nlf.d.pomp`. `logLik` applied to such an object returns the log quasi likelihood. The `$` method allows extraction of arbitrary slots from the `nlf.d.pomp` object.

Author(s)

Stephen P. Ellner, Bruce E. Kendall, Aaron A. King

References

The following papers describe and motivate the NLF approach to model fitting:

Ellner, S. P., Bailey, B. A., Bobashev, G. V., Gallant, A. R., Grenfell, B. T. and Nychka D. W. (1998) Noise and nonlinearity in measles epidemics: combining mechanistic and statistical approaches to population modeling. *American Naturalist* **151**, 425–440.

Kendall, B. E., Briggs, C. J., Murdoch, W. W., Turchin, P., Ellner, S. P., McCauley, E., Nisbet, R. M. and Wood S. N. (1999) Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches. *Ecology* **80**, 1789–1805.

Kendall, B. E., Ellner, S. P., McCauley, E., Wood, S. N., Briggs, C. J., Murdoch, W. W. and Turchin, P. (2005) Population cycles in the pine looper moth (*Bupalus piniarius*): dynamical tests of mechanistic hypotheses. *Ecological Monographs* **75**, 259–276.

ou2

Two-dimensional discrete-time Ornstein-Uhlenbeck process

Description

ou2 is a pomp object encoding a bivariate discrete-time Ornstein-Uhlenbeck process.

Details

If the state process is $X(t) = (x_1(t), x_2(t))$, then

$$X(t+1) = \alpha X(t) + \sigma \epsilon(t),$$

where α and σ are 2x2 matrices, σ is lower-triangular, and $\epsilon(t)$ is standard bivariate normal. The observation process is $Y(t) = (y_1(t), y_2(t))$, where $y_i(t) \sim \text{normal}(x_i(t), \tau)$. The functions `rprocess`, `dprocess`, `rmeasure`, `dmeasure`, and `skeleton` are implemented using compiled C code for computational speed: see the source code for details.

See Also

[pomp](#)

Examples

```
pompExample(ou2)
plot(ou2)
coef(ou2)
x <- simulate(ou2)
plot(x)
pf <- pfilter(ou2, Np=1000)
logLik(pf)
```

parmat	<i>Create a matrix of parameters</i>
--------	--------------------------------------

Description

`parmat` is a utility that makes a vector of parameters suitable for use in **pomp** functions.

Usage

```
parmat(params, nrep = 1)
```

Arguments

<code>params</code>	named numeric vector or matrix of parameters.
<code>nrep</code>	number of replicates (columns) desired.

Value

`parmat` returns a matrix consisting of `nrep` copies of `params`.

Author(s)

Aaron A. King

Examples

```
## generate a bifurcation diagram for the Ricker map
pompExample(ricker)
p <- parmat(coef(ricker), nrep=500)
p["r",] <- exp(seq(from=1.5, to=4, length=500))
x <- trajectory(ricker, times=seq(from=1000, to=2000, by=1), params=p)
matplot(p["r",], x["N",,], pch=., col=black, xlab="log(r)", ylab="N", log=x)
```

Particle filter

Particle filter

Description

A plain vanilla sequential Monte Carlo (particle filter) algorithm. Resampling is performed at each observation.

Usage

```
## S4 method for signature pomp
pfilter(object, params, Np, tol = 1e-17,
        max.fail = Inf, pred.mean = FALSE, pred.var = FALSE,
        filter.mean = FALSE, filter.traj = FALSE, save.states = FALSE,
        save.params = FALSE, verbose = getOption("verbose"), ...)
## S4 method for signature pfilterd.pomp
pfilter(object, params, Np, tol, ...)
## S4 method for signature pfilterd.pomp
logLik(object, ...)
## S4 method for signature pfilterd.pomp
cond.logLik(object, ...)
## S4 method for signature pfilterd.pomp
eff.sample.size(object, ...)
## S4 method for signature pfilterd.pomp
pred.mean(object, pars, ...)
## S4 method for signature pfilterd.pomp
pred.var(object, pars, ...)
## S4 method for signature pfilterd.pomp
filter.mean(object, pars, ...)
## S4 method for signature pfilterd.pomp
filter.traj(object, vars, ...)
```

Arguments

object	An object of class pomp or inheriting class pomp.
params	A npars x Np numeric matrix containing the parameters corresponding to the initial state values in xstart. This must have a 'rownames' attribute. If it desired that all particles should share the same parameter values, one one may supply params as a named numeric vector.

<code>Np</code>	<p>the number of particles to use. This may be specified as a single positive integer, in which case the same number of particles will be used at each timestep. Alternatively, if one wishes the number of particles to vary across timesteps, one may specify <code>Np</code> either as a vector of positive integers of length <code>length(time(object, t0=TRUE))</code></p> <p>or as a function taking a positive integer argument. In the latter case, <code>Np(k)</code> must be a single positive integer, representing the number of particles to be used at the k-th timestep: <code>Np(0)</code> is the number of particles to use going from <code>timezero(object)</code> to <code>time(object)[1]</code>, <code>Np(1)</code>, from <code>timezero(object)</code> to <code>time(object)[1]</code>, and so on, while when <code>T=length(time(object, t0=TRUE))</code>, <code>Np(T)</code> is the number of particles to sample at the end of the time-series. When <code>object</code> is of class <code>mif</code>, this is by default the same number of particles used in the <code>mif</code> iterations.</p>
<code>tol</code>	positive numeric scalar; particles with likelihood less than <code>tol</code> are considered to be “lost”. A filtering failure occurs when, at some time point, all particles are lost. When all particles are lost, the conditional likelihood at that time point is set to <code>tol</code> .
<code>max.fail</code>	integer; the maximum number of filtering failures allowed. If the number of filtering failures exceeds this number, execution will terminate with an error. By default, <code>max.fail</code> is set to infinity, so no error can be triggered.
<code>pred.mean</code>	logical; if <code>TRUE</code> , the prediction means are calculated for the state variables and parameters.
<code>pred.var</code>	logical; if <code>TRUE</code> , the prediction variances are calculated for the state variables and parameters.
<code>filter.mean</code>	logical; if <code>TRUE</code> , the filtering means are calculated for the state variables and parameters.
<code>filter.traj</code>	logical; if <code>TRUE</code> , a filtered trajectory is returned for the state variables and parameters.
<code>save.states, save.params</code>	logical. If <code>save.states=TRUE</code> , the state-vector for each particle at each time is saved in the <code>saved.states</code> slot of the returned <code>pfilterd.pomp</code> object. If <code>save.params=TRUE</code> , the parameter-vector for each particle at each time is saved in the <code>saved.params</code> slot of the returned <code>pfilterd.pomp</code> object.
<code>verbose</code>	logical; if <code>TRUE</code> , progress information is reported as <code>pfilter</code> works.
<code>pars</code>	Names of parameters.
<code>vars</code>	Names of state variables.
<code>...</code>	additional arguments that override the defaults.

Value

An object of class `pfilterd.pomp`. This class inherits from class `pomp` and contains the following additional slots:

pred.mean, pred.var, filter.mean matrices of prediction means, variances, and filter means, respectively. In each of these, the rows correspond to states and parameters (if appropriate), in that order, the columns to successive observations in the time series contained in `object`.

filter.traj filtered trajectories returned by the particle filter

eff.sample.size numeric vector containing the effective number of particles at each time point.

cond.loglik numeric vector containing the conditional log likelihoods at each time point.

saved.states If `pfilter` was called with `save.states=TRUE`, this is the list of state-vectors at each time point, for each particle. It is a length-`ntimes` list of `nvars`-by-`Np` arrays. In particular, `saved.states[[t]][,i]` can be considered a sample from $f[X_t|y_{1:t}]$.

saved.params If `pfilter` was called with `save.params=TRUE`, this is the list of parameter-vectors at each time point, for each particle. It is a length-`ntimes` list of `npars`-by-`Np` arrays. In particular, `saved.params[[t]][,i]` is the parameter portion of the i -th particle at time t .

Np, tol, nfail the number of particles used, failure tolerance, and number of filtering failures, respectively.

loglik the estimated log-likelihood.

These can be accessed using the `$` operator as if the returned object were a list. Note that if the argument `params` is a named vector, then these parameters are included in the `params` slot of the returned `pfilterd.pomp` object.

Methods

logLik Extracts the estimated log likelihood.

cond.logLik Extracts the estimated conditional log likelihood

$$\ell_t(\theta) = \text{Prob}[y_t|y_1, \dots, y_{t-1}],$$

where y_t are the data, at time t .

eff.sample.size Extracts the (time-dependent) estimated effective sample size, computed as

$$\left(\sum_i w_{it}^2 \right)^{-1},$$

where w_{it} is the normalized weight of particle i at time t .

pred.mean, pred.var Extract the mean and variance of the approximate prediction distribution. This prediction distribution is that of

$$X_t|y_1, \dots, y_{t-1},$$

where X_t, y_t are the state vector and data, respectively, at time t .

filter.mean Extract the mean of the filtering distribution, which is that of

$$X_t|y_1, \dots, y_t,$$

where X_t, y_t are the state vector and data, respectively, at time t .

Author(s)

Aaron A. King

References

M. S. Arulampalam, S. Maskell, N. Gordon, & T. Clapp. A Tutorial on Particle Filters for Online Nonlinear, Non-Gaussian Bayesian Tracking. *IEEE Trans. Sig. Proc.* 50:174–188, 2002.

See Also

[pomp](#), [mif](#), [pmcmc](#), [bsmc2](#), and the tutorials on the [package website](#).

Examples

```
pompExample(gompertz)
pf <- pfilter(gompertz,Np=1000) ## use 1000 particles
plot(pf)
logLik(pf)
cond.logLik(pf) ## conditional log-likelihoods
eff.sample.size(pf) ## effective sample size
logLik(pfilter(pf)) ## run it again with 1000 particles
## run it again with 2000 particles
pf <- pfilter(pf,Np=2000,filter.mean=TRUE)
fm <- filter.mean(pf) ## extract the filtering means
```

Particle Markov Chain Monte Carlo

The particle Markov chain Metropolis-Hastings algorithm

Description

The Particle MCMC algorithm for estimating the parameters of a partially-observed Markov process. Running `pmcmc` causes a particle random-walk Metropolis-Hastings Markov chain algorithm to run for the specified number of proposals.

Usage

```
## S4 method for signature pomp
pmcmc(object, Nmcmc = 1, start, proposal, Np,
      tol = 1e-17, max.fail = Inf, verbose = getOption("verbose"), ...)
## S4 method for signature pfilterd.pomp
pmcmc(object, Nmcmc = 1, Np, tol, ...)
## S4 method for signature pmcmc
pmcmc(object, Nmcmc, start, proposal, Np, tol,
      max.fail = Inf, verbose = getOption("verbose"), ...)
## S4 method for signature pmcmc
continue(object, Nmcmc = 1, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>Nmcmc</code>	The number of PMCMC iterations to perform.
<code>start</code>	named numeric vector; the starting guess of the parameters.
<code>proposal</code>	optional function that draws from the proposal distribution. Currently, the proposal distribution must be symmetric for proper inference: it is the user's responsibility to ensure that it is. Several functions that construct appropriate proposal function are provided: see MCMC proposal functions for more information.
<code>Np</code>	a positive integer; the number of particles to use in each filtering operation.
<code>tol</code>	numeric scalar; particles with log likelihood below <code>tol</code> are considered to be "lost". A filtering failure occurs when, at some time point, all particles are lost.
<code>max.fail</code>	integer; maximum number of filtering failures permitted. If the number of failures exceeds this number, execution will terminate with an error.
<code>verbose</code>	logical; if TRUE, print progress reports.
<code>...</code>	additional arguments that override the defaults.

Value

An object of class `pmcmc`.

Re-running PMCMC Iterations

To re-run a sequence of PMCMC iterations, one can use the `pmcmc` method on a `pmcmc` object. By default, the same parameters used for the original PMCMC run are re-used (except for `tol`, `max.fail`, and `verbose`, the defaults of which are shown above). If one does specify additional arguments, these will override the defaults.

Continuing PMCMC Iterations

One can continue a series of PMCMC iterations from where one left off using the `continue` method. A call to `pmcmc` to perform `Nmcmc=m` iterations followed by a call to `continue` to perform `Nmcmc=n` iterations will produce precisely the same effect as a single call to `pmcmc` to perform `Nmcmc=m+n` iterations. By default, all the algorithmic parameters are the same as used in the original call to `pmcmc`. Additional arguments will override the defaults.

Details

`pmcmc` implements an MCMC algorithm in which the true likelihood of the data is replaced by an unbiased estimate computed by a particle filter. This gives an asymptotically correct Bayesian procedure for parameter estimation (Andrieu and Roberts, 2009). An extension to give a correct Bayesian posterior distribution of unobserved state variables (as in Andrieu et al, 2010) has not yet been implemented.

Methods

`conv.rec(object, pars)` returns the columns of the convergence-record matrix corresponding to the names in `pars` as an object of class `mcmc` or `mcmc.list`.

`filter.traj(object, vars)` returns filter trajectories.

plot Diagnostic plots.

logLik Returns the value in the `loglik` slot.

c Concatenates `pmcmc` objects into a `pmcmcList`.

`covmat(object, start, thin, expand)` computes the empirical covariance matrix of the MCMC samples beginning with iteration `start` and thinning by factor `thin`. It expands this by a factor expand^2/n , where `n` is the number of parameters estimated. By default, `expand=2.38`. The intention is that the resulting matrix is a suitable input to the proposal function `mvn.rw`.

Author(s)

Edward L. Ionides, Aaron A. King, Sebastian Funk

References

C. Andrieu, A. Doucet and R. Holenstein, Particle Markov chain Monte Carlo methods, J. R. Stat. Soc. B, to appear, 2010.

C. Andrieu and G.O. Roberts, The pseudo-marginal approach for efficient computation, Ann. Stat. 37:697-725, 2009.

See Also

[pomp](#), [pfilter](#), [MCMC proposal distributions](#), and the tutorials on the [package website](#).

Examples

```
## Not run:
library(pomp)

pompExample(ou2)

pmcmc(
  pomp(ou2, dprior=Csnippet("
lik = dnorm(alpha_2, -0.5, 1, 1) + dnorm(alpha_3, 0.3, 1, 1);
lik = (give_log) ? lik : exp(lik);"),
  paramnames=c("alpha.2", "alpha.3")),
  Nmcmc=2000, Np=500, verbose=TRUE,
  proposal=mvn.rw.adaptive(rw.sd=c(alpha.2=0.01, alpha.3=0.01),
    scale.start=200, shape.start=100)) -> chain
continue(chain, Nmcmc=2000, proposal=mvn.rw(covmat(chain))) -> chain
plot(chain)
chain <- pmcmc(chain)
plot(chain)

require(coda)
trace <- window(conv.rec(chain, c("alpha.2", "alpha.3")), start=2000)
```

```

rejectionRate(trace)
effectiveSize(trace)
autocorr.diag(trace)

plot(trace)
heidel.diag(trace)
geweke.diag(trace)

## End(Not run)

```

pomp constructor

Constructor of the basic POMP object

Description

This function constructs a pomp object, encoding a partially-observed Markov process model together with a uni- or multivariate time series. One implements the model by specifying its *components*, each of which can be written as R functions or, for much greater computational efficiency, using C code. The preferred way to specify most components (as detailed below) is through the use of [Csnippets](#), snippets of C that are compiled and linked into a running R session.

Usage

```

## S4 method for signature data.frame
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model,
      skeleton, skeleton.type = c("map", "vectorfield"), skelmap.delta.t = 1,
      initializer, rprior, dprior, params, covar, tcovar,
      obsnames, statenames, paramnames, covarnames, zeronames,
      PACKAGE, fromEstimationScale, toEstimationScale, globals)

## S4 method for signature numeric
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model,
      skeleton, skeleton.type = c("map", "vectorfield"), skelmap.delta.t = 1,
      initializer, rprior, dprior, params, covar, tcovar,
      obsnames, statenames, paramnames, covarnames, zeronames,
      PACKAGE, fromEstimationScale, toEstimationScale, globals)

## S4 method for signature matrix
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model,
      skeleton, skeleton.type = c("map", "vectorfield"), skelmap.delta.t = 1,
      initializer, rprior, dprior, params, covar, tcovar,
      obsnames, statenames, paramnames, covarnames, zeronames,
      PACKAGE, fromEstimationScale, toEstimationScale, globals)

## S4 method for signature pomp
pomp(data, times, t0, ..., rprocess, dprocess, rmeasure, dmeasure,
      measurement.model, skeleton, skeleton.type, skelmap.delta.t,
      initializer, rprior, dprior, params, covar, tcovar,

```

```
obsnames, statenames, paramnames, covarnames, zeronames,
PACKAGE, fromEstimationScale, toEstimationScale, globals)
```

Arguments

- | | |
|-------------|--|
| data, times | <p>The time series data and times at which observations are made. data can be specified as a vector, a matrix, a data-frame. Alternatively, a pomp object can be supplied in the data argument.</p> <p>If data is a numeric vector, times must be a numeric vector of the same length. If data is a matrix, it should have dimensions nobs x ntimes, where nobs is the number of observed variables and ntimes is the number of times at which observations were made (i.e., each column is a distinct observation of the nobs variables). In this case, times must be given as a numeric vector (of length ntimes).</p> <p>If data is a data-frame, times must name the column of observation times. Note that, in this case, data will be internally coerced to an array with storage-mode 'double'.</p> <p>times must be numeric and strictly increasing.</p> |
| t0 | <p>The zero-time, at which the stochastic dynamical system is to be initialized. This must be no later than the time of the first observation, i.e., $t0 \leq \text{times}[1]$.</p> |
| rprocess | <p>optional function; a function of prototype</p> <pre>rprocess(xstart, times, params, ...)</pre> <p>that simulates from the unobserved process. The form of this function is given below. pomp provides a number of plugins that construct appropriate rprocess arguments corresponding to several common stochastic simulation algorithms. See below for more details.</p> |
| dprocess | <p>optional function; a function of prototype</p> <pre>dprocess(x, times, params, log, ...)</pre> <p>that evaluates the likelihood of a sequence of consecutive state transitions. The form of this function is given below. It is not typically necessary (or even feasible) to define dprocess. This functionality is provided to support future algorithm development.</p> |
| rmeasure | <p>optional; the measurement model simulator. This can be specified in one of four ways:</p> <ol style="list-style-type: none"> 1. as a function of prototype <pre>rmeasure(x, t, params, ...)</pre> <p>that makes a draw from the observation process given states x, time t, and parameters params.</p> 2. as the name of a native (compiled) routine of type <pre>pomp_measure_model_simulator</pre> <p>as defined in the header file 'pomp.h'. (To view the header file, execute <code>file.show(system.file("include/pomp.h", package="pomp"))</code> in an R session.)</p> |

3. using the formula-based `measurement.model` facility (see below).
4. as a snippet of C code (via [Csnippet](#)) that draws from the observation process as above. The last is typically the preferred option, as it results in much faster code execution.

`dmeasure` optional; the measurement model probability density function. This can be specified in one of four ways:

1. as a function of prototype
`dmeasure(y,x,t,params,log,...)`
 that computes the p.d.f. of y given x , t , and $params$.
2. as the name of a native (compiled) routine of type
`pomp_measure_model_density`
 as defined in the header file 'pomp.h'. (To view the header file, execute `file.show(system.file("include/pomp.h",package="pomp"))` in an R session.)
3. using the formula-based `measurement.model` facility (see below).
4. as a snippet of C code (via [Csnippet](#)) that computes the p.d.f. as above.

The last is typically the preferred option, as it results in much faster code execution. As might be expected, if `log=TRUE`, this function should return the log likelihood.

`measurement.model`

optional; a formula or list of formulae, specifying the measurement model. These formulae are parsed internally to generate `rmeasure` and `dmeasure` functions. If `measurement.model` is given it overrides any specification of `rmeasure` or `dmeasure`. **NB:** This is a convenience function, primarily designed to facilitate exploration; it will typically be possible to accelerate measurement model computations by writing `dmeasure` and/or `rmeasure` using [Csnippets](#).

`skeleton`, `skeleton.type`, `skelmap.delta.t`

The function `skeleton` specifies the deterministic skeleton of the unobserved Markov process. If we are dealing with a discrete-time Markov process, its deterministic skeleton is a map: indicate this by specifying `skeleton.type="map"`. In this case, the default assumption is that time advances 1 unit per iteration of the map; to change this, set `skelmap.delta.t` to the appropriate time-step. If we are dealing with a continuous-time Markov process, its deterministic skeleton is a vectorfield: indicate this by specifying `skeleton.type="vectorfield"`. The skeleton function can be specified in one of three ways:

1. as an R function of prototype
`skeleton(x,t,params,...)`
 that evaluates the deterministic skeleton at state x and time t given the parameters $params$,
2. as the name of a native (compiled) routine of type
`pomp_skeleton`
 as defined in the header file 'pomp.h'. (To view the header file, execute `file.show(system.file("include/pomp.h",package="pomp"))` in an R session.)

	<ol style="list-style-type: none"> 3. as a snippet of C code (via Csnippet) that performs this evaluation. The latter is typically the preferred option, for reasons of computational efficiency.
initializer	<p>The initializer gives the parameterization of the initial state of the unobserved Markov process. Specifically, given a vector of parameters, <code>params</code> and an initial time, <code>t0</code>, the initializer determines the state vector at time <code>t0</code>.</p> <p>By default, any parameters in <code>params</code>, the names of which end in “.0”, are assumed to be initial values of states. To initialize the unobserved state process, these are simply copied over as initial conditions. The names of the resulting state variables are obtained by dropping the “.0” suffix.</p> <p>A custom initializer can be supplied here in one of two formats:</p> <ol style="list-style-type: none"> 1. as an R function of prototype <pre>initializer(params,t0,...)</pre> that yields initial conditions for the state process when given a vector, <code>params</code>, of parameters. 2. as a snippet of C code (via Csnippet). See the Csnippet help for rules on writing a custom initializer.
rprior	<p>optional; function drawing a sample from a prior distribution on parameters. This can be specified in one of three ways:</p> <ol style="list-style-type: none"> 1. as an R function of prototype <pre>rprior(params,...)</pre> that makes a draw from the prior distribution given <code>params</code>, 2. as the name of a native (compiled) routine of type <pre>pomp_rprior</pre> as defined in the header file ‘<code>pomp.h</code>’, or (To view the header file, execute <code>file.show(system.file("include/pomp.h",package="pomp"))</code> in an R session.) 3. as a snippet of C code (via Csnippet). <p>As above, the latter is typically preferable.</p>
dprior	<p>optional; function evaluating the prior distribution. This can be specified in one of three ways:</p> <ol style="list-style-type: none"> 1. as an R function of prototype <pre>dprior(params,log=FALSE,...)</pre> that evaluates the prior probability density, 2. as the name of a native (compiled) routine of type <pre>pomp_dprior</pre> as defined in the header file ‘<code>pomp.h</code>’, or (To view the header file, execute <code>file.show(system.file("include/pomp.h",package="pomp"))</code> in an R session.) 3. as a snippet of C code (via Csnippet). <p>As above, the latter is typically preferable.</p>
params	<p>optional named numeric vector of parameters. This will be coerced internally to storage mode double.</p>

covar, tcovar	<p>An optional matrix or data frame of covariates: covar is the table of covariates (one column per variable); tcovar the corresponding times (one entry per row of covar).</p> <p>covar can be specified as either a matrix or a data frame. In either case the columns are taken to be distinct covariates. If covar is a data frame, tcovar can be either the name or the index of the time variable.</p> <p>If a covariate table is supplied, then the value of each of the covariates is interpolated as needed. The resulting interpolated values are passed to the corresponding functions as a numeric vector named covars; see below.</p>
obsnames, statenames, paramnames, covarnames	<p>Optional character vectors specifying the names of observables, state variables, parameters, and covariates, respectively. These are only used in the event that one or more of the basic functions (rprocess, dprocess, rmeasure, dmeasure, skeleton, rprior, dprior) are defined using Csnippet or native routines.</p>
zernames	<p>optional character vector specifying the names of accumulator variables (see below).</p>
PACKAGE	<p>An optional string giving the name of the dynamically loaded library in which any native routines are to be found.</p>
fromEstimationScale, toEstimationScale	<p>Optional functions specifying parameter transformations. Many algorithms for parameter estimation search an unconstrained space of parameters. When working with such an algorithm and a model for which the parameters are constrained, it can be useful to transform parameters. toEstimationScale and fromEstimationScale are transformations from the model scale to the estimation scale, and vice versa, respectively. These functions must have arguments params and toEstimationScale should transform parameters from the scale that rprocess, dprocess, rmeasure, dmeasure, and skeleton use internally to the scale used in estimation. fromEstimationScale should be the inverse of toEstimationScale. The parameter transformations can be defined (as above) using either R functions, native routines, or Csnippets.</p> <p>Note that it is the user's responsibility to make sure that these transformations are mutually inverse. If obj is the constructed pomp object, and coef(obj) is non-empty, a simple check of this property is</p> <pre> x <- coef(obj, transform=TRUE) obj1 <- obj coef(obj1, transform=TRUE) <- x identical(coef(obj), coef(obj1)) identical(coef(obj1, transform=TRUE), x). </pre> <p>By default, both functions are the identity transformation. See the demos, <code>demo(package="pomp")</code>, pompExample, and the tutorials on the package website for examples.</p>
globals	<p>optional character; C code that will be included in the source for (and therefore hard-coded into) the shared-object library created when the call to pomp uses Csnippets. If no Csnippets are used, globals has no effect.</p>

... Any additional arguments given to `pomp` will be stored in the `pomp` object and passed as arguments to each of the basic functions whenever they are evaluated.

Value

`pomp` returns an object of class `pomp`. If `data` is an object of class `pomp`, then by default the returned `pomp` object is identical to `data`. If additional arguments are given, these override the defaults.

Important note

It is not typically necessary (or even feasible) to define all of the components `rprocess`, `dprocess`, `rmeasure`, `dmeasure`, and `skeleton` in any given problem. Each algorithm makes use of only a subset of these components. Any algorithm requiring a component that has not been defined will return an informative error.

The state process model

Specification of process-model codes `rprocess` and/or `dprocess` in most cases is facilitated by **pomp's** so-called [plugins](#), which have been developed to handle common use-cases. Currently, if one's process model evolves in discrete time or one is willing to make such an approximation (e.g., via an Euler approximation), then the [euler.sim](#), [discrete.time.sim](#), and [onestep.sim](#) plugins for `rprocess` and [onestep.dens](#) plugin for `dprocess` are available. In addition, for exact simulation of certain continuous-time Markov chains, an implementation of Gillespie's algorithm is available (see [gillespie.sim](#)). To learn more about the use of plugins, consult the help documentation ([plugins](#)) and the tutorials on the [package website](#). Several of the demos and examples make use of these as well.

In specific cases, it may be possible to obtain increased computational efficiency by writing custom versions of `rprocess` and/or `dprocess` instead of using the plugins. If such custom versions are desired, the following describes how these functions should be written.

rprocess If the plugins are not used `rprocess` must be an R function with at least the following arguments: `xstart`, `times`, `params`, and `...`. It can also take additional arguments. It is guaranteed that these will be filled with the corresponding elements the user has included as additional arguments in the construction of the `pomp` object.

In calls to `rprocess`, `xstart` can be assumed to be an `nvar` x `nrep` matrix; its rows correspond to components of the state vector and columns correspond to independent realizations of the process. `params` will similarly be an `npar` x `nrep` matrix with rows corresponding to parameters and columns corresponding to independent realizations. Note that the columns of `params` correspond to those of `xstart`; in particular, they will agree in number. Both `xstart` and `params` are guaranteed to have rownames.

`rprocess` must return a rank-3 array with rownames. Suppose `x` is the array returned. Then `dim(x)=c(nvars,nrep,ntimes)`, where `ntimes` is the length of the vector `times`. `x[,j,k]` is the value of the state process in the `j`-th realization at time `times[k]`. In particular, `x[, ,1]` must be identical to `xstart`. The rownames of `x` must correspond to those of `xstart`.

dprocess If the plugins are not used, `dprocess` must have at least the following arguments: `x`, `times`, `params`, `log`, and `...`. It may take additional arguments: again, these will be filled with the corresponding elements the user defines when the `pomp` object is constructed.

In calls to `dprocess`, `x` may be assumed to be an `nvars` x `nrep` x `ntimes` array, where these terms have the same meanings as above. `params` will be a matrix with rows corresponding to

individual parameters and columns corresponding to independent realizations. The columns of `params` correspond to those of `x`; in particular, they will agree in number. Both `x` and `params` are guaranteed to have rownames.

`dprocess` must return a `nrep x ntimes-1` matrix. Suppose `d` is the array returned. `d[j,k]` is the probability density of the transition from state `x[,j,k-1]` at time `times[k-1]` to state `x[,j,k]` at time `times[k]`. If `log=TRUE`, then the log of the pdf must be returned.

In writing this function, you may assume that the transitions are consecutive. It should be clear that, but for this assumption, it will in general be impossible to write the transition probabilities explicitly. In such cases, algorithms that make no use of `dprocess`, which are said to have the “plug and play” property, are useful. Most of the algorithms in **pomp** have this property. In particular, **at present, no methods in pomp make use of dprocess.**

The observation process model

The following is a guide to writing the measurement model components as R functions. For a description on how to write these components using [Csnippets](#), see the tutorials on the [package website](#).

`rmeasure` if provided, must take at least the arguments `x`, `t`, `params`, and `...`. It may take additional arguments, which will be filled with user-specified data as above. `x` will be a named numeric vector of length `nvars` (which has the same meaning as above). `t` will be a scalar quantity, the time at which the measurement is made. `params` will be a named numeric vector of length `nparams`. The `rmeasure` function may take additional arguments which will be filled with user-specified data as above.

`rmeasure` must return a named numeric vector of length `nobs`, the number of observable variables.

`dmeasure` if provided, must take at least the arguments `y`, `x`, `t`, `params`, `log`, and `...`. `y` will be a named numeric vector of length `nobs` containing (actual or simulated) values of the observed variables; `x` will be a named numeric vector of length `nvar` containing state variables; `params` will be a named numeric vector containing parameters; and `t` will be a scalar, the corresponding observation time. The `dmeasure` function may take additional arguments which will be filled with user-specified data as above.

`dmeasure` must return a single numeric value, the probability density of `y` given `x` at time `t`. If `log=TRUE`, then `dmeasure` should return the log of the probability density.

The deterministic skeleton

The following describes how to specify the deterministic skeleton as an R function. For a description on how to write this component using [Csnippets](#), see the tutorials on the [package website](#) and the [Csnippet](#) help.

If `skeleton` if provided, must have at least the arguments `x`, `t`, `params`, and `...`. `x` is a numeric vector containing the coordinates of a point in state space at which evaluation of the skeleton is desired. `t` is a numeric value giving the time at which evaluation of the skeleton is desired. Of course, these will be irrelevant in the case of an autonomous skeleton. `params` is a numeric vector holding the parameters. `skeleton` may take additional arguments, which will be filled, as above, with user-specified data.

`skeleton` must return a numeric vector of the same length as `x`, which contains the value `vectorfield` (if the dynamical system is continuous) or the value of the map (if the dynamical system is discrete), at the point `x` at time `t`.

The state-process initializer

if provided, must have at least the arguments `params`, `t0`, and `...` `params` will be a named numeric vector of parameters. `t0` will be the time at which initial conditions are desired. `initializer` must return a named numeric vector of initial states.

Covariates

If the `pomp` object contains covariates (via the `covar` argument; see above), then whenever any of the R functions described above are called, they will each be supplied with an additional argument `covars`. This will be a named numeric vector containing the (interpolated) values of the covariates at the time `t`. In particular, `covars` will have one value for each column of the covariate table.

Accumulator variables

In formulating models, one often wishes to define a state variable that will accumulate some quantity over the interval between successive observations. **pomp** provides a facility to make such features more convenient. Specifically, variables named in the `pomp`'s `zeronames` argument will be set to zero immediately following each observation. See `euler.sir` and the tutorials on the [package website](#) for examples.

Warning

Some error checking is done by `pomp`, but complete error checking is impossible. If the user-specified functions do not conform to the above specifications, then the results may be invalid. In particular, if both `rmeasure` and `dmeasure` are specified, the user should verify that these two functions correspond to the same probability distribution. If `skeleton` is specified, the user is responsible for verifying that it corresponds to a deterministic skeleton of the model. Each **pomp**-package algorithm uses some subset of the five basic functions (`rprocess`, `dprocess`, `rmeasure`, `dmeasure`, `skeleton`). If an algorithm requires a component that has not been specified, an informative error will be generated.

Author(s)

Aaron A. King

See Also

[pomp methods](#), [pomp low-level interface](#), [process model plugins](#)

Examples

```
## Not run:
pompExample()
pomp.home <- system.file("examples", package="pomp")
pomp.examples <- list.files(pomp.home)
```

```

file.show(
  file.path(pomp.home, pomp.examples),
  header=paste("=====", pomp.examples, "=====")
)

## End(Not run)

```

pomp methods

*Functions for manipulating, displaying, and extracting information
from objects of the pomp class*

Description

This page documents the various methods that allow one to extract information from, display, plot, and modify pomp objects.

Usage

```

## S4 method for signature pomp
coef(object, pars, transform = FALSE, ...)
## S4 replacement method for signature pomp
coef(object, pars, transform = FALSE, ...) <- value
## S4 method for signature pomp
obs(object, vars, ...)
## S4 method for signature pomp
partrans(object, params, dir = c("fromEstimationScale",
  "toEstimationScale", "forward", "inverse"), ...)
## S4 method for signature pomp
plot(x, y, variables, panel = lines,
  nc = NULL, yax.flip = FALSE,
  mar = c(0, 5.1, 0, if (yax.flip) 5.1 else 2.1),
  oma = c(6, 0, 5, 0), axes = TRUE, ...)
## S4 method for signature pomp
print(x, ...)
## S4 method for signature pomp
show(object)
## S4 method for signature pomp
states(object, vars, ...)
## S4 method for signature pomp
time(x, t0 = FALSE, ...)
## S4 replacement method for signature pomp
time(object, t0 = FALSE, ...) <- value
## S4 method for signature pomp
timezero(object, ...)
## S4 replacement method for signature pomp
timezero(object, ...) <- value
## S4 method for signature pomp

```

```

window(x, start, end, ...)
## S4 method for signature pomp
as(object, class)

```

Arguments

object, x	The pomp object.
pars	optional character; names of parameters to be retrieved or set.
vars	optional character; names of observed variables to be retrieved.
transform	optional logical; should the parameter transformations be applied?
value	numeric; values to be assigned.
params	a vector or matrix of parameters to be transformed.
dir	direction of the transformation. <code>dir="forward"</code> applies the transformation from the “natural” scale to the “internal” scale. This is the transformation specified by the parameter <code>.transform</code> argument to <code>pomp</code> ; it is stored in the <code>‘par.trans’</code> slot of object. <code>dir="inverse"</code> applies the inverse transformation (stored in the <code>‘par.untrans’</code> slot).
t0	logical; if TRUE on a call to <code>time</code> , the zero time is prepended to the time vector; if TRUE on a call to <code>time<-</code> , the first element in <code>value</code> is taken to be the initial time.
start, end	start and end times of the window.
class	character; name of the class to which object should be coerced.
y	ignored.
variables	optional character; names of variables to plot.
panel	a function of prototype <code>panel(x, col, bg, pch, type, ...)</code> which gives the action to be carried out in each panel of the display.
nc	the number of columns to use. Defaults to 1 for up to 4 series, otherwise to 2.
yax.flip	logical; if TRUE, the y-axis (ticks and numbering) should flip from side 2 (left) to 4 (right) from series to series.
mar, oma	the ‘par’ settings for ‘mar’ and ‘oma’ to use. Modify with care!
axes	logical; indicates if x- and y- axes should be drawn.
...	Further arguments (either ignored or passed to underlying functions).

Details

coef `coef(object)` returns the contents of the `params` slot of object. `coef(object,pars)` returns only those parameters named in `pars`.

```
coef(object,transform=TRUE)
```

returns

```
parameter.inv.transform(coef(object))
```

, where `parameter.inv.transform` is the user parameter inverse transformation function specified when object was created. Likewise,

```
coef(object,pars,transform=TRUE)

returns

parameter.inv.transform(coef(object))[pars]
```

coef<- Assigns values to the params slot of the pomp object. `coef(object) <- value` has the effect of replacing the parameters of object with value. If `coef(object)` exists, then `coef(object,pars) <- value` replaces those parameters of object named in pars with the elements of value; the names of value are ignored. If some of the names in pars do not already name parameters in `coef(object)`, then they are concatenated. If `coef(object)` does not exist, then `coef(object,pars) <- value` assigns value to the parameters of object; in this case, the names of object will be pars and the names of value will be ignored. `coef(object,transform=TRUE) <- value` assigns `parameter.transform(value)` to the params slot of object. Here, `parameter.transform` is the parameter transformation function specified when object was created. `coef(object,pars,transform=TRUE) <- value` first, discards any names the value may have, sets `names(value) <- pars`, and then replaces the elements of object's params slot `parameter.transform(value)`. In this case, if some of the names in pars do not already name parameters in `coef(object,transform=TRUE)`, then they are concatenated.

obs `obs(object)` returns the array of observations. `obs(object,vars)` gives just the observations of variables named in vars. vars may specify the variables by position or by name.

states `states(object)` returns the array of states. `states(object,vars)` gives just the state variables named in vars. vars may specify the variables by position or by name.

time `time(object)` returns the vector of observation times. `time(object,t0=TRUE)` returns the vector of observation times with the zero-time t0 prepended.

time<- `time(object) <- value` replaces the observation times slot (times) of object with value. `time(object,t0=TRUE) <- value` has the same effect, but the first element in value is taken to be the initial time. The second and subsequent elements of value are taken to be the observation times. Those data and states (if they exist) corresponding to the new times are retained.

timezero, timezero<- `timezero(object)` returns the zero-time t0. `timezero(object) <- value` sets the zero-time to value.

window `window(x,start=t1,end=t2)` returns a new pomp object, identical to x but with only the data in the window between times t1 and t2 (inclusive). By default, start is the time of the first observation and end is the time of the last.

show Displays the pomp object.

plot Plots the data and state trajectories (if the latter exist). Additional arguments are passed to the low-level plotting routine.

print Prints the pomp object in a nice way.

as A pomp object can be coerced to a data frame via

```
as(object,"data.frame").
```

The data frame contains the times, the data, and the state trajectories, if they exist.

Author(s)

Aaron A. King

See Also[pomp](#), [pomp low-level interface](#), [simulate](#), [pfilter](#), [probe](#).

POMP simulation

*Simulations of a partially-observed Markov process***Description**`simulate` generates simulations of the state and measurement processes.**Usage**

```
## S4 method for signature pomp
simulate(object, nsim = 1, seed = NULL, params,
        states = FALSE, obs = FALSE, times, t0,
        as.data.frame = FALSE, include.data = FALSE, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>nsim</code>	The number of simulations to perform. Note that the number of replicates will be <code>nsim</code> times <code>ncol(xstart)</code> .
<code>seed</code>	optional; if set, the pseudorandom number generator (RNG) will be initialized with <code>seed</code> . the random seed to use. The RNG will be restored to its original state afterward.
<code>params</code>	either a named numeric vector or a numeric matrix with rownames. The parameters to use in simulating the model. If <code>params</code> is not given, then the contents of the <code>params</code> slot of <code>object</code> will be used, if they exist.
<code>states</code>	Do we want the state trajectories?
<code>obs</code>	Do we want data-frames of the simulated observations?
<code>times, t0</code>	<code>times</code> specifies the times at which simulated observations will be made. <code>t0</code> specifies the start time (the time at which the initial conditions hold). The default for <code>times</code> is <code>times=time(object,t0=FALSE)</code> and <code>t0=timezero(object)</code> , respectively.
<code>as.data.frame, include.data</code>	logical; if <code>as.data.frame=TRUE</code> , the results are returned as a data-frame. A factor variable, 'sim', distinguishes one simulation from another. If, in addition, <code>include.data=TRUE</code> , the original data are included as an additional 'simulation'. If <code>as.data.frame=FALSE</code> , <code>include.data</code> is ignored.
<code>...</code>	further arguments that are currently ignored.

Details

Simulation of the state process and of the measurement process are each accomplished by a single call to the user-supplied `rprocess` and `rmeasure` functions, respectively. This makes it possible for the user to write highly optimized code for these potentially expensive computations.

Value

If `states=FALSE` and `obs=FALSE` (the default), a list of `nsim` pomp objects is returned. Each has a simulated data set, together with the parameters used (in slot `params`) and the state trajectories also (in slot `states`). If `times` is specified, then the simulated observations will be at times `times`.

If `nsim=1`, then a single pomp object is returned (and not a singleton list).

If `states=TRUE` and `obs=FALSE`, simulated state trajectories are returned as a rank-3 array with dimensions `nvar` x `(ncol(params)*nsim)` x `ntimes`. Here, `nvar` is the number of state variables and `ntimes` the length of the argument `times`. The measurement process is not simulated in this case.

If `states=FALSE` and `obs=TRUE`, simulated observations are returned as a rank-3 array with dimensions `nobs` x `(ncol(params)*nsim)` x `ntimes`. Here, `nobs` is the number of observables.

If both `states=TRUE` and `obs=TRUE`, then a named list is returned. It contains the state trajectories and simulated observations as above.

Author(s)

Aaron A. King

See Also

[pomp](#)

Examples

```
pompExample(ou2)
x <- simulate(ou2, seed=3495485, nsim=10)
x <- simulate(ou2, seed=3495485, nsim=10, states=TRUE, obs=TRUE)
x <- simulate(ou2, seed=3495485, nsim=10, obs=TRUE,
              as.data.frame=TRUE, include.data=TRUE)
```

Description

`spect` estimates the power spectrum of time series data and model simulations and compares the results. It can be used to diagnose goodness of fit and/or as the basis for frequency-domain parameter estimation (`spect.match`).

`spect.match` tries to match the power spectrum of the model to that of the data. It calls an optimizer to adjust model parameters to minimize the discrepancy between simulated and actual data.

Usage

```
## S4 method for signature pomp
spect(object, params, vars, kernel.width, nsim, seed = NULL,
      transform = identity,
      detrend = c("none", "mean", "linear", "quadratic"),
      ...)
## S4 method for signature spect.pomp
spect(object, params, vars, kernel.width, nsim, seed = NULL, transform,
      detrend, ...)
spect.match(object, start, est = character(0),
            vars, nsim, seed = NULL,
            kernel.width, transform = identity,
            detrend = c("none", "mean", "linear", "quadratic"),
            weights, method = c("subplex", "Nelder-Mead", "SANN"),
            verbose = getOption("verbose"),
            eval.only = FALSE, fail.value = NA, ...)
```

Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>params</code>	optional named numeric vector of model parameters. By default, <code>params=coef(object)</code> .
<code>vars</code>	optional; names of observed variables for which the power spectrum will be computed. This must be a subset of <code>rownames(obs(object))</code> . By default, the spectrum will be computed for all observables.
<code>kernel.width</code>	width parameter for the smoothing kernel used for calculating the estimate of the spectrum.
<code>nsim</code>	number of model simulations to be computed.
<code>seed</code>	optional; if non-NULL, the random number generator will be initialized with this seed for simulations. See simulate .
<code>transform</code>	function; this transformation will be applied to the observables prior to estimation of the spectrum, and prior to any detrending.
<code>detrend</code>	de-trending operation to perform. Options include no detrending, and subtraction of constant, linear, and quadratic trends from the data. Detrending is applied to each data series and to each model simulation independently.
<code>weights</code>	optional. The mismatch between model and data is measured by a weighted average of mismatch at each frequency. By default, all frequencies are weighted equally. <code>weights</code> can be specified either as a vector (which must have length

	equal to the number of frequencies) or as a function of frequency. If the latter, <code>weights(freq)</code> must return a nonnegative weight for each frequency.
<code>start</code>	named numeric vector; the initial guess of parameters.
<code>est</code>	character vector; the names of parameters to be estimated.
<code>method</code>	Optimization method. Choices are <code>subplex</code> and any of the methods used by <code>optim</code> .
<code>verbose</code>	logical; print diagnostic messages?
<code>eval.only</code>	logical; if TRUE, no optimization is attempted. Instead, the probe-mismatch value is simply evaluated at the <code>start</code> parameters.
<code>fail.value</code>	optional scalar; if non-NA, this value is substituted for non-finite values of the objective function.
<code>...</code>	Additional arguments. In the case of <code>spect</code> , these are currently ignored. In the case of <code>spect.match</code> , these are passed to <code>optim</code> or <code>subplex</code> in the control list.

Details

A call to `spect` results in the estimation of the power spectrum for the (transformed, detrended) data and `nsim` model simulations. The results of these computations are stored in an object of class `spect.pomp`.

A call to `spect.match` results in an attempt to optimize the agreement between model and data spectrum over the parameters named in `est`. The results, including coefficients of the fitted model and power spectra of fitted model and data, are stored in an object of class `spect.matched.pomp`.

Value

`spect` returns an object of class `spect.pomp`, which is derived from class `pomp` and therefore has all the slots of that class. In addition, `spect.pomp` objects have the following slots:

kernel.width width parameter of the smoothing kernel used.

transform transformation function used.

freq numeric vector of the frequencies at which the power spectrum is estimated.

datSPEC, simSPEC estimated power spectra for data and simulations, respectively.

pvals one-sided p-values: fraction of the simulated spectra that differ more from the mean simulated spectrum than does the data. The metric used is L^2 distance.

detrend detrending option used.

`spect.match` returns an object of class `spect.matched.pomp`, which is derived from class `spect.pomp` and therefore has all the slots of that class. In addition, `spect.matched.pomp` objects have the following slots:

est, weights, fail.value values of the corresponding arguments in the call to `spect.match`.

evals number of function and gradient evaluations by the optimizer. See `optim`.

value Value of the objective function.

convergence, msg Convergence code and message from the optimizer. See `optim`.

Author(s)

Daniel C. Reuman, Cai GoGwilt, Aaron A. King

References

D.C. Reuman, R.A. Desharnais, R.F. Costantino, O. Ahmad, J.E. Cohen (2006) Power spectra reveal the influence of stochasticity on nonlinear population dynamics. *Proceedings of the National Academy of Sciences* **103**, 18860-18865.

D.C. Reuman, R.F. Costantino, R.A. Desharnais, J.E. Cohen (2008) Color of environmental noise affects the nonlinear dynamics of cycling, stage-structured populations. *Ecology Letters*, **11**, 820-830.

See Also

[pomp](#), [probe](#)

Examples

```
pompExample(ou2)
good <- spect(
  ou2,
  vars=c("y1", "y2"),
  kernel.width=3,
  detrend="mean",
  nsim=500
)
summary(good)
plot(good)

ou2.bad <- ou2
coef(ou2.bad, c("x1.0", "x2.0", "alpha.1", "alpha.4")) <- c(0, 0, 0.1, 0.2)
bad <- spect(
  ou2.bad,
  vars=c("y1", "y2"),
  kernel.width=3,
  detrend="mean",
  nsim=500
)
summary(bad)
plot(bad)
```

Description

Several simple and configurable probes are provided with in the package. These can be used directly and as templates for custom probes.

Usage

```

probe.mean(var, trim = 0, transform = identity, na.rm = TRUE)
probe.median(var, na.rm = TRUE)
probe.var(var, transform = identity, na.rm = TRUE)
probe.sd(var, transform = identity, na.rm = TRUE)
probe.marginal(var, ref, order = 3, diff = 1, transform = identity)
probe.nlar(var, lags, powers, transform = identity)
probe.acf(var, lags, type = c("covariance", "correlation"),
          transform = identity)
probe.ccf(vars, lags, type = c("covariance", "correlation"),
          transform = identity)
probe.period(var, kernel.width, transform = identity)
probe.quantile(var, prob, transform = identity)

```

Arguments

<code>var, vars</code>	character; the name(s) of the observed variable(s).
<code>trim</code>	the fraction of observations to be trimmed (see mean).
<code>transform</code>	transformation to be applied to the data before the probe is computed.
<code>na.rm</code>	if TRUE, remove all NA observations prior to computing the probe.
<code>kernel.width</code>	width of modified Daniell smoothing kernel to be used in power-spectrum computation: see kernel .
<code>prob</code>	a single probability; the quantile to compute: see quantile .
<code>lags</code>	In <code>probe.ccf</code> , a vector of lags between time series. Positive lags correspond to x advanced relative to y ; negative lags, to the reverse. In <code>probe.nlar</code> , a vector of lags present in the nonlinear autoregressive model that will be fit to the actual and simulated data. See Details, below, for a precise description.
<code>powers</code>	the powers of each term (corresponding to lags) in the the nonlinear autoregressive model that will be fit to the actual and simulated data. See Details, below, for a precise description.
<code>type</code>	Compute autocorrelation or autocovariance?
<code>ref</code>	empirical reference distribution. Simulated data will be regressed against the values of <code>ref</code> , sorted and, optionally, differenced. The resulting regression coefficients capture information about the shape of the marginal distribution. A good choice for <code>ref</code> is the data itself.
<code>order</code>	order of polynomial regression.
<code>diff</code>	order of differencing to perform.
<code>...</code>	Additional arguments to be passed through to the probe computation.

Details

Each of these functions is relatively simple. See the source code for a complete understanding of what each does.

`probe.mean`, `probe.median`, `probe.var`, `probe.sd` return functions that compute the mean, median, variance, and standard deviation of variable `var`, respectively.

`probe.period` returns a function that estimates the period of the Fourier component of the `var` series with largest power.

`probe.marginal` returns a function that regresses the marginal distribution of variable `var` against the reference distribution `ref`. If `diff>0`, the data and the reference distribution are first differenced `diff` times and centered. Polynomial regression of order `order` is used. This probe returns order regression coefficients (the intercept is zero).

`probe.nlar` returns a function that fit a nonlinear (polynomial) autoregressive model to the univariate series (variable `var`). Specifically, a model of the form $y_t = \sum \beta_k y_{t-\tau_k}^{p_k} + \epsilon_t$ will be fit, where τ_k are the lags and p_k are the powers. The data are first centered. This function returns the regression coefficients, β_k .

`probe.acf` returns a function that, if `type=="covariance"`, computes the autocovariance of variable `var` at lags `lags`; if `type=="correlation"`, computes the autocorrelation of variable `var` at lags `lags`.

`probe.ccf` returns a function that, if `type=="covariance"`, computes the cross covariance of the two variables named in `vars` at lags `lags`; if `type=="correlation"`, computes the cross correlation.

`probe.quantile` returns a function that estimates the `prob`-th quantile of variable `var`.

Value

A call to any one of these functions returns a probe function, suitable for use in `probe` or `probe.match`. That is, the function returned by each of these takes a data array (such as comes from a call to `obs`) as input and returns a single numerical value.

Author(s)

Daniel C. Reuman (d.reuman at imperial dot ac dot uk)

Aaron A. King (kingaa at umich dot edu)

References

B. E. Kendall, C. J. Briggs, W. M. Murdoch, P. Turchin, S. P. Ellner, E. McCauley, R. M. Nisbet, S. N. Wood Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches, *Ecology*, 80:1789–1805, 1999.

S. N. Wood Statistical inference for noisy nonlinear ecological dynamic systems, *Nature*, 466: 1102–1104, 2010.

See Also

[pomp](#)

 Probes and synthetic likelihood

Probe a partially-observed Markov process by computing summary statistics and the synthetic likelihood.

Description

probe applies one or more “probes” to time series data and model simulations and compares the results. It can be used to diagnose goodness of fit and/or as the basis for “probe-matching”, a generalized method-of-moments approach to parameter estimation. probe.match calls an optimizer to adjust model parameters to do probe-matching, i.e., to minimize the discrepancy between simulated and actual data. This discrepancy is measured using the “synthetic likelihood” as defined by Wood (2010). probe.match.objfun constructs an objective function for probe-matching suitable for use in optim-like optimizers.

Usage

```
## S4 method for signature pomp
probe(object, probes, params, nsim, seed = NULL, ...)
## S4 method for signature probed.pomp
probe(object, probes, params, nsim, seed, ...)
## S4 method for signature pomp
probe.match.objfun(object, params, est, probes, nsim,
  seed = NULL, fail.value = NA, transform = FALSE, ...)
## S4 method for signature probed.pomp
probe.match.objfun(object, probes, nsim, seed, ...)
## S4 method for signature pomp
probe.match(object, start, est = character(0),
  probes, nsim, seed = NULL,
  method = c("subplex", "Nelder-Mead", "SANN", "BFGS",
    "sannbox", "nloptr"),
  verbose = getOption("verbose"),
  fail.value = NA, transform = FALSE, ...)
## S4 method for signature probed.pomp
probe.match(object, probes, nsim, seed,
  ..., verbose = getOption("verbose"))
## S4 method for signature probe.matched.pomp
probe.match(object, est, probes,
  nsim, seed, transform, fail.value, ...,
  verbose = getOption("verbose"))
## S4 method for signature probed.pomp
logLik(object, ...)
## S4 method for signature probed.pomp
values(object, ...)
```


Arguments

<code>object</code>	An object of class <code>pomp</code> .
<code>probes</code>	A single probe or a list of one or more probes. A probe is simply a scalar- or vector-valued function of one argument that can be applied to the data array of a <code>pomp</code> . A vector-valued probe must always return a vector of the same size. A number of useful examples are provided with the package: see probe functions).
<code>params</code>	optional named numeric vector of model parameters. By default, <code>params=coef(object)</code> .
<code>nsim</code>	The number of model simulations to be computed.
<code>seed</code>	optional; if non-NULL, the random number generator will be initialized with this seed for simulations. See simulate-pomp .
<code>start</code>	named numeric vector; the initial guess of parameters.
<code>est</code>	character vector; the names of parameters to be estimated.
<code>method</code>	Optimization method. Choices refer to algorithms used in optim , subplex , and nloptr .
<code>verbose</code>	logical; print diagnostic messages?
<code>fail.value</code>	optional numeric scalar; if non-NA, this value is substituted for non-finite values of the objective function. It should be a large number (i.e., bigger than any legitimate values the objective function is likely to take).
<code>transform</code>	logical; if TRUE, optimization is performed on the transformed scale.
<code>...</code>	Additional arguments. In the case of <code>probe</code> , these are currently ignored. In the case of <code>probe.match</code> , these are passed to the optimizer (one of optim , subplex , nloptr , or sannbox). These are passed via the optimizer's control list (in the case of <code>optim</code> , <code>subplex</code> , and <code>sannbox</code>) or the <code>opts</code> list (in the case of <code>nloptr</code>).

Details

A call to `probe` results in the evaluation of the probe(s) in `probes` on the data. Additionally, `nsim` simulated data sets are generated (via a call to [simulate](#)) and the probe(s) are applied to each of these. The results of the probe computations on real and simulated data are stored in an object of class `probed.pomp`.

A call to `probe.match` results in an attempt to optimize the agreement between model and data, as measured by the specified probes, over the parameters named in `est`. The results, including coefficients of the fitted model and values of the probes for data and fitted-model simulations, are stored in an object of class [probe.matched.pomp](#).

The objective function minimized by `probe.match` — in a form suitable for use with `optim`-like optimizers — is created by a call to `probe.match.objfun`. Specifically, `probe.match.objfun` will return a function that takes a single numeric-vector argument that is assumed to contain the parameters named in `est`, in that order. This function will return the negative synthetic log likelihood for the probes specified.

Value

`probe` returns an object of class `probed.pomp`. `probed.pomp` is derived from the [pomp](#) class and therefore have all the slots of `pomp`. In addition, a `probed.pomp` class has the following slots:

probes list of the probes applied.

datvals, simvals values of each of the probes applied to the real and simulated data, respectively.

quantiles fraction of simulations with probe values less than the value of the probe of the data.

pvals two-sided p-values: fraction of the `simvals` that deviate more extremely from the mean of the `simvals` than does `datvals`.

synth.loglik the log synthetic likelihood (Wood 2010). This is the likelihood assuming that the probes are multivariate-normally distributed.

`probe.match` returns an object of class `probe.matched.pomp`, which is derived from class `probed.pomp`. `probe.matched.pomp` objects therefore have all the slots above plus the following:

est, transform, fail.value values of the corresponding arguments in the call to `probe.match`.

value value of the objective function at the optimum.

evals number of function and gradient evaluations by the optimizer. See [optim](#).

convergence, msg Convergence code and message from the optimizer. See [optim](#) and [nloptr](#).

`probe.match.objfun` returns a function suitable for use as an objective function in an [optim](#)-like optimizer.

Methods

plot displays diagnostic plots.

summary displays summary information.

values extracts the realized values of the probes on the data and on the simulations as a data frame in long format. The variable `.id` indicates whether the probes are from the data or simulations.

logLik returns the synthetic likelihood for the probes. NB: in general, this is not the same as the likelihood.

as, as.data.frame when a ‘`probed.pomp`’ is coerced to a ‘`data.frame`’, the first row gives the probes applied to the data; the rest of the rows give the probes evaluated on simulated data. The rownames of the result can be used to distinguish these.

In addition, slots of this object can be accessed via the `$` operator.

Author(s)

Daniel C. Reuman, Aaron A. King

References

B. E. Kendall, C. J. Briggs, W. M. Murdoch, P. Turchin, S. P. Ellner, E. McCauley, R. M. Nisbet, S. N. Wood Why do populations cycle? A synthesis of statistical and mechanistic modeling approaches, *Ecology*, 80:1789–1805, 1999.

S. N. Wood Statistical inference for noisy nonlinear ecological dynamic systems, *Nature*, 466: 1102–1104, 2010.

See Also

[pomp](#), [probe functions](#), [spect](#), and the tutorials on the [package website](#).

Examples

```

pompExample(ou2)
good <- probe(
  ou2,
  probes=list(
    y1.mean=probe.mean(var="y1"),
    y2.mean=probe.mean(var="y2"),
    y1.sd=probe.sd(var="y1"),
    y2.sd=probe.sd(var="y2"),
    extra=function(x)max(x["y1",])
  ),
  nsim=500
)
summary(good)
plot(good)

bad <- probe(
  ou2,
  params=c(alpha.1=0.1,alpha.4=0.2,x1.0=0,x2.0=0,
    alpha.2=-0.5,alpha.3=0.3,
    sigma.1=3,sigma.2=-0.5,sigma.3=2,
    tau=1),
  probes=list(
    y1.mean=probe.mean(var="y1"),
    y2.mean=probe.mean(var="y2"),
    y1.sd=probe.sd(var="y1"),
    y2.sd=probe.sd(var="y2"),
    extra=function(x)range(x["y1",])
  ),
  nsim=500
)
summary(bad)
plot(bad)

```

Process model plugins *Plug-ins for state-process models*

Description

The functions documented here can be used to specify the `rprocess` and `dprocess` slots for a `pomp` model. There are options for discrete- and continuous-time Markov processes.

Usage

```

onestep.sim(step.fun, PACKAGE)
euler.sim(step.fun, delta.t, PACKAGE)
discrete.time.sim(step.fun, delta.t = 1, PACKAGE)
gillespie.sim(rate.fun, v, d, PACKAGE)
onestep.dens(dens.fun, PACKAGE)

```

Arguments

<code>step.fun</code>	<p>This can be either an R function, the name of a compiled, dynamically loaded native function containing the model simulator, or a Csnippet. The latter is the preferred option. It should be written to take a single Euler step from a single point in state space.</p> <p>For an explanation and examples on the use of Csnippet to write fast simulators, see the tutorials on the package website.</p> <p>If it is an R function, it should have prototype</p> <pre>step.fun(x,t,params,delta.t,...).</pre> <p>Here, <code>x</code> is a named numeric vector containing the value of the state process at time <code>t</code>, <code>params</code> is a named numeric vector containing parameters, and <code>delta.t</code> is the length of the Euler time-step.</p> <p>If <code>step.fun</code> is the name of a native function, it must be of type</p> <pre>pomp_onestep_sim</pre> <p>as defined in the header file 'pomp.h', which is included with the pomp package. Do</p> <pre>file.show(system.file("include/pomp.h",package="pomp"))</pre> <p>to view this header file. For details on how to write such codes, see Details.</p>
<code>rate.fun</code>	<p>This can be either an R function, a Csnippet, or the name of a compiled, dynamically loaded native function that computes the transition rates. If it is an R function, it should be of the form <code>rate.fun(j,x,t,params,...)</code>. Here, <code>j</code> is the number of the event, <code>x</code> is a named numeric vector containing the value of the state process at time <code>t</code> and <code>params</code> is a named numeric vector containing parameters.</p> <p>For examples on the use of Csnippet to write fast simulators easily, see tutorials on the package website.</p> <p>If <code>rate.fun</code> is a native function, it must be of type</p> <pre>pomp_ssa_rate_fn</pre> <p>as defined in the header 'pomp.h', which is included with the package. For details on how to write such codes, see Details.</p>
<code>v, d</code>	<p>Matrices that specify the continuous-time Markov process in terms of its elementary events. Each should have dimensions <code>nvar</code> x <code>nevent</code>, where <code>nvar</code> is the number of state variables and <code>nevent</code> is the number of elementary events. <code>v</code> describes the changes that occur in each elementary event: it will usually comprise the values 1, -1, and 0 according to whether a state variable is incremented, decremented, or unchanged in an elementary event. <code>d</code> is a binary matrix that describes the dependencies of elementary event rates on state variables: <code>d[i,j]</code> will have value 1 if event rate <code>j</code> must be updated as a result of a change in state variable <code>i</code> and 0 otherwise</p>
<code>dens.fun</code>	<p>This can be either an R function, a Csnippet, or a compiled, dynamically loaded native function containing the model transition log probability density function. If it is an R function, it should be of the form</p>

```
dens.fun(x1,x2,t1,t2,params,...).
```

Here, `x1` and `x2` are named numeric vectors containing the values of the state process at times `t1` and `t2`, `params` is a named numeric vector containing parameters.

If `dens.fun` is the name of a native function, it should be of type

```
pomp_onestep_pdf
```

as defined in the header ‘`pomp.h`’, which is included with the **pomp** package. This function should return the log likelihood of a transition from `x1` at time `t1` to `x2` at time `t2`, assuming that no intervening transitions have occurred. For details on how to write such codes, see Details.

`delta.t` Size of Euler time-steps.

`PACKAGE` an optional argument that specifies to which dynamically loaded library we restrict the search for the native routines. If this is “base”, we search in the R executable itself. This argument is ignored if `step.fun`, `rate.fn`, or `dens.fun` is provided as an R function or a [Csnippet](#).

Details

`onestep.sim` is the appropriate choice when it is possible to simulate the change in state from one time to another, regardless of how large the interval between them is. To use `onestep.sim`, you must write a function `step.fun` that will advance the state process from one arbitrary time to another. `euler.sim` is appropriate when one cannot do this but can compute the change in state via a sequence of smaller steps. This is desirable, for example, if one is simulating a continuous time process but is willing to approximate it using an Euler approach. `discrete.time.sim` is appropriate when the process evolves in discrete time.

To use `euler.sim` or `discrete.time.sim`, you must write a function `step.fun` that will take a single Euler step, of size at most `delta.t`. The functions `euler.sim` and `discrete.time.sim` will create simulators that take as many steps as needed to get from one time to another. See below for information on how `euler.sim` chooses the actual step size it uses.

`gillespie.sim` allows exact simulation of a continuous-time, discrete-state Markov process using Gillespie’s algorithm. This is an “event-driven” approach: correspondingly, to use `gillespie.sim`, you must write a function `rate.fun` that computes the rates of each elementary kind of event and specify two matrices (`d`, `v`) that describe, respectively, the dependencies of each rate and the consequences of each event.

`onestep.dens` will generate a suitable `dprocess` function when one can compute the likelihood of a given state transition simply by knowing the states at two times under the assumption that the state has not changed between the times. This is typically possible, for instance, when the `rprocess` function is implemented using `onestep.sim`, `euler.sim`, or `discrete.time.sim`. [NB: currently, there are no high-level algorithms in **pomp** that use `dprocess`. This function is provided for completeness only, with an eye toward future development.]

If `step.fun` is written as an R function, it must have at least the arguments `x`, `t`, `params`, `delta.t`, and `...`. On a call to this function, `x` will be a named vector of state variables, `t` a scalar time, and `params` a named vector of parameters. The length of the Euler step will be `delta.t`. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t`, of the covariates. This is

accomplished via interpolation of the user-supplied covariate table. Additional arguments may be given: these will be filled by the correspondingly-named elements in the `userdata` slot of the `pomp` object (see [pomp](#)).

If `step.fun` is written in a native language, it must be a function of type

```
pomp_onestep_sim
```

as specified in the header ‘`pomp.h`’ included with the package. Execute

```
file.show(system.file("include/pomp.h", package="pomp.h"))
```

to view this file.

If `rate.fun` is written as an R function, it must have at least the arguments `j`, `x`, `t`, `params`, and `...`. Here, `j` is the an integer that indicates for which of the elementary events the current rate is desired. `x` is a named vector containing the value of the state process at time `t`, and `params` is a named vector containing parameters. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t`, of the covariates. This is accomplished via interpolation of the covariate table. If `rate.fun` is a native function, it must be of type

```
pomp_ssa_rate_fn
```

as defined in the header ‘`pomp.h`’; see above for instructions on how to view this file.

In writing `dens.fun`, you must assume that no state transitions have occurred between `t1` and `t2`. If `dens.fun` is written as an R function, it must have at least the arguments `x1`, `x2`, `t1`, `t2`, `params`, and `...`. On a call to this function, `x1` and `x2` will be named vectors of state variables at times `t1` and `t2`, respectively. The named vector `params` contains the parameters. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t1`, of the covariates. If the argument `covars` is included and a covariate table has been included in the `pomp` object, then on a call to this function, `covars` will be filled with the values, at time `t1`, of the covariates. This is accomplished via interpolation of the covariate table. As above, any additional arguments will be filled by the correspondingly-named elements in the `userdata` slot of the `pomp` object (see [pomp](#)). If `dens.fun` is written in a native language, it must be a function of type

```
pomp_onestep_pdf
```

as defined in the header ‘`pomp.h`’ included with the package; see above for instructions on how to view this file.

Value

`onestep.sim`, `euler.sim`, `discrete.time.sim`, and `gillespie.sim` each return functions suitable for use as the argument `rprocess` argument in [pomp](#).

`onestep.dens` returns a function suitable for use as the argument `dprocess` in [pomp](#).

Author(s)

Aaron A. King

See Also

[pomp](#) and the tutorials on the [package website](#).

ricker

Ricker model with Poisson observations.

Description

ricker is a pomp object encoding a stochastic Ricker model with Poisson measurement error.

Details

The state process is $N_{t+1} = rN_t \exp(-N_t + e_t)$, where the e_t are i.i.d. normal random deviates with zero mean and variance σ^2 . The observed variables y_t are distributed as $\text{Poisson}(\phi N_t)$.

See Also

[pomp](#), [gompertz](#), and the tutorials on the [package website](#).

Examples

```
pompExample(ricker)
plot(ricker)
coef(ricker)
```

rw2

Two-dimensional random-walk process

Description

rw2 is a pomp object encoding a 2-D normal random walk.

Details

The random-walk process is fully but noisily observed.

See Also

[pomp](#), [ou2](#)

Examples

```
pompExample(rw2)
plot(rw2)
x <- simulate(rw2, nsim=10, seed=20348585L, params=c(x1.0=0, x2.0=0, s1=1, s2=3, tau=1))
plot(x[[1]])
```

Simulated annealing *Simulated annealing with box constraints.*

Description

sannbox is a straightforward implementation of simulated annealing with box constraints.

Usage

```
sannbox(par, fn, control = list(), ...)
```

Arguments

<code>par</code>	Initial values for the parameters to be optimized over.
<code>fn</code>	A function to be minimized, with first argument the vector of parameters over which minimization is to take place. It should return a scalar result.
<code>control</code>	A named list of control parameters. See ‘Details’.
<code>...</code>	ignored.

Details

The control argument is a list that can supply any of the following components:

`trace` Non-negative integer. If positive, tracing information on the progress of the optimization is produced. Higher values may produce more tracing information.

`fnscale` An overall scaling to be applied to the value of `fn` during optimization. If negative, turns the problem into a maximization problem. Optimization is performed on `fn(par)/fnscale`.

`parscale` A vector of scaling values for the parameters. Optimization is performed on `par/parscale` and these should be comparable in the sense that a unit change in any element produces about a unit change in the scaled value.

`maxit` The total number of function evaluations: there is no other stopping criterion. Defaults to 10000.

`temp` starting temperature for the cooling schedule. Defaults to 1.

`tmax` number of function evaluations at each temperature. Defaults to 10.

`candidate.dist` function to randomly select a new candidate parameter vector. This should be a function with three arguments, the first being the current parameter vector, the second the temperature, and the third the parameter scaling. By default, `candidate.dist` is

```
function(par, temp, scale) rnorm(n=length(par), mean=par, sd=scale*temp)
```

.

`sched` cooling schedule. A function of a three arguments giving the temperature as a function of iteration number and the control parameters `temp` and `tmax`. By default, `sched` is

```
function(k, temp, tmax) temp/log(((k-1)/tmax)*tmax+exp(1))
```

.

Alternatively, one can supply a numeric vector of temperatures. This must be of length at least `maxit`.

Value

sannbox returns a list with components:

counts two-element integer vector. The first number gives the number of calls made to `fn`. The second number is provided for compatibility with `optim` and will always be NA.

convergence provided for compatibility with `optim`; will always be 0.

final.params last tried value of `par`.

final.value value of `fn` corresponding to `final.params`.

par best tried value of `par`.

value value of `fn` corresponding to `par`.

Author(s)

Daniel Reuman, Aaron A. King

See Also

[traj.match](#), [probe.match](#).

sir

Compartmental epidemiological models

Description

`euler.sir` is a `pomp` object encoding a simple seasonal SIR model. Simulation is performed using an Euler multinomial approximation. `gillespie.sir` has the same model implemented using Gillespie's algorithm. `bbs` is a nonseasonal SIR model together with data from a 1978 outbreak of influenza in a British boarding school.

Details

This and similar examples are discussed and constructed in tutorials available on the [package website](#).

The boarding school influenza outbreak is described in Anonymous (1978).

References

Anonymous (1978). Influenza in a boarding school. *British Medical Journal* 1:587

See Also

[pomp](#) and the tutorials on the [package website](#).

Examples

```
pompExample(euler.sir)
plot(euler.sir)
plot(simulate(euler.sir,seed=20348585))
coef(euler.sir)

pompExample(gillespie.sir)
plot(gillespie.sir)
plot(simulate(gillespie.sir,seed=20348585))
coef(gillespie.sir)

pompExample(bbs)
plot(bbs)
coef(bbs)
```

Trajectory matching	<i>Parameter estimation by fitting the trajectory of a model's deterministic skeleton to data</i>
---------------------	---

Description

This function attempts to match trajectories of a model's deterministic skeleton to data. Trajectory matching is equivalent to maximum likelihood estimation under the assumption that process noise is entirely absent, i.e., that all stochasticity is measurement error. Accordingly, this method uses only the skeleton and dmeasure components of a POMP model.

Usage

```
## S4 method for signature pomp
traj.match(object, start, est = character(0),
           method = c("Nelder-Mead", "subplex", "SANN", "BFGS",
                     "sannbox", "nloptr"),
           transform = FALSE, ...)
## S4 method for signature traj.matched.pomp
traj.match(object, est, transform, ...)
## S4 method for signature pomp
traj.match.objfun(object, params, est, transform = FALSE, ...)
```

Arguments

object	A pomp object. If object has no skeleton slot, an error will be generated.
start	named numeric vector containing an initial guess for parameters. By default start=coef(object) if the latter exists.
params	optional named numeric vector of parameters. This should contain all parameters needed by the skeleton and dmeasure slots of object. In particular, any parameters that are to be treated as fixed should be present here. Parameter values given in params for parameters named in est will be ignored. By default, params=coef(object) if the latter exists.

<code>est</code>	character vector containing the names of parameters to be estimated. In the case of <code>traj.match.objfun</code> , the objective function that is constructed will assume that its argument contains the parameters in this order.
<code>method</code>	Optimization method. Choices are <code>subplex</code> , “ <code>sannbox</code> ”, and any of the methods used by <code>optim</code> .
<code>transform</code>	logical; if TRUE, optimization is performed on the transformed scale.
<code>...</code>	Extra arguments that will be passed either to the optimizer (<code>optim</code> , <code>subplex</code> , <code>nloptr</code> , or <code>sannbox</code> , via their control (<code>optim</code> , <code>subplex</code> , <code>sannbox</code>) or <code>opts</code> (<code>nloptr</code>) lists) or to the ODE integrator. In <code>traj.match</code> , extra arguments will be passed to the optimizer. In <code>traj.match.objfun</code> , extra arguments are passed to <code>trajectory</code> . If extra arguments are needed by both optimizer and <code>trajectory</code> , construct an objective function first using <code>traj.match.objfun</code> , then give this objective function to the optimizer.

Details

In **pomp**, trajectory matching is the term used for maximizing the likelihood of the data under the assumption that there is no process noise. Specifically, `traj.match` calls an optimizer (`optim`, `subplex`, and `sannbox` are the currently supported options) to minimize an objective function. For any value of the model parameters, this objective function is calculated by

1. computing the deterministic trajectory of the model given the parameters. This is the trajectory returned by `trajectory`, which relies on the model’s deterministic skeleton as specified in the construction of the `pomp` object object.
2. evaluating the negative log likelihood of the data under the measurement model given the deterministic trajectory and the model parameters. This is accomplished via the model’s `dmeasure` slot. The negative log likelihood is the objective function’s value.

The objective function itself — in a form suitable for use with `optim`-like optimizers — is created by a call to `traj.match.objfun`. Specifically, `traj.match.objfun` will return a function that takes a single numeric-vector argument that is assumed to contain the parameters named in `est`, in that order.

Value

`traj.match` returns an object of class `traj.matched.pomp`. This class inherits from class `pomp` and contains the following additional slots:

transform, est the values of these arguments on the call to `traj.match`.

evals number of function and gradient evaluations by the optimizer. See `optim`.

value value of the objective function. Larger values indicate better fit (i.e., `traj.match` attempts to maximize this quantity).

convergence, msg convergence code and message from the optimizer. See `optim`.

Available methods for objects of this type include `summary` and `logLik`. The other slots of this object can be accessed via the `$` operator.

`traj.match.objfun` returns a function suitable for use as an objective function in an `optim`-like optimizer.

See Also

[trajectory](#), [pomp](#), [optim](#), [subplex](#)

Examples

```
pompExample(ou2)
true.p <- c(
  alpha.1=0.9,alpha.2=0,alpha.3=-0.4,alpha.4=0.99,
  sigma.1=2,sigma.2=0.1,sigma.3=2,
  tau=1,
  x1.0=50,x2.0=-50
)
simdata <- simulate(ou2,nsim=1,params=true.p,seed=43553)
guess.p <- true.p
res <- traj.match(
  simdata,
  start=guess.p,
  est=c(alpha.1,alpha.3,alpha.4,x1.0,x2.0,tau),
  maxit=2000,
  method="Nelder-Mead",
  reltol=1e-8
)

summary(res)

plot(range(time(res)),range(c(obs(res),states(res))),type="n",xlab="time",ylab="x,y")
points(y1~time,data=as(res,"data.frame"),col=blue)
points(y2~time,data=as(res,"data.frame"),col=red)
lines(x1~time,data=as(res,"data.frame"),col=blue)
lines(x2~time,data=as(res,"data.frame"),col=red)

pompExample(ricker)
ofun <- traj.match.objfun(ricker,est=c("r","phi"),transform=TRUE)
optim(fn=ofun,par=c(2,0),method="BFGS")

pompExample(bbs)
## some options are passed to the ODE integrator
ofun <- traj.match.objfun(bbs,est=c("beta","gamma"),transform=TRUE,hmax=0.001,rtol=1e-6)
optim(fn=ofun,par=c(0,-1),method="Nelder-Mead",control=list(reltol=1e-10))
```

Utilities for reproducibility

Tools for reproducible computations.

Description

On cooking shows, recipes requiring lengthy baking or stewing are prepared beforehand. The `bake` and `stew` functions perform analogously: an R computation is performed and stored in a named file. If the function is called again and the file is present, the computation is not executed; rather, the

results are loaded from the file in which they were previously stored. Moreover, via their optional seed argument, `bake` and `stew` can control the pseudorandom-number generator (RNG) for greater reproducibility. After the computation is finished, these functions restore the pre-existing RNG state to avoid side effects.

The `freeze` function doesn't save results, but does set the RNG state to the specified value and restore it after the computation is complete.

Usage

```
bake(file, expr, seed, kind = NULL, normal.kind = NULL)
stew(file, expr, seed, kind = NULL, normal.kind = NULL)
freeze(expr, seed, kind = NULL, normal.kind = NULL)
```

Arguments

<code>file</code>	Name of the binary data file in which the result will be stored or retrieved, as appropriate. For <code>bake</code> , this will contain a single R object and hence be an RDS file (extension 'rds'); for <code>stew</code> , this will contain one or more named R objects and hence be an RDA file (extension 'rda').
<code>expr</code>	Expression to be evaluated.
<code>seed, kind, normal.kind</code>	Optional. To set the state and, optionally, kind of RNG used. See set.seed .

Details

Both `bake` and `stew` first test to see whether file exists. If it does, `bake` reads it using [readRDS](#) and returns the resulting object. By contrast, `stew` loads the file using [load](#) and copies the objects it contains into the user's workspace (or the environment of the call to `stew`).

If file does not exist, then both `bake` and `stew` evaluate the expression `expr`; they differ in the results that they save. `bake` saves the value of the evaluated expression to file as a single R object. The name of that object is not saved. By contrast, `stew` creates a local environment within which `expr` is evaluated; all objects in that environment are saved (by name) in file.

Value

`bake` returns the value of the evaluated expression `expr`. Other objects created in the evaluation of `expr` are discarded along with the temporary, local environment created for the evaluation.

The latter behavior differs from that of `stew`, which returns the names of the objects created during the evaluation of `expr`. After `stew` completes, these objects exist in the parent environment (that from which `stew` was called).

`freeze` returns the value of evaluated expression `expr`. However, `freeze` evaluates `expr` within the parent environment, so other objects created in the evaluation of `expr` will therefore exist after `freeze` completes.

Author(s)

Aaron A. King

Examples

```
## Not run:
bake(file="example1.rds",{
  x <- runif(1000)
  mean(x)
})

stew(file="example2.rds",{
  x <- runif(10)
  y <- rnorm(n=10,mean=3*x+5,sd=2)
})

plot(x,y)

## End(Not run)

freeze(runif(3),seed=5886730)
freeze(runif(3),seed=5886730)
```

Index

- *Topic **datasets**
 - blowflies, [10](#)
 - Childhood disease incidence data, [11](#)
 - dacca, [16](#)
 - Example pomp models, [20](#)
 - gompertz, [22](#)
 - ou2, [38](#)
 - pomp-package, [3](#)
 - ricker, [71](#)
 - rw2, [71](#)
 - sir, [73](#)
- *Topic **design**
 - design, [17](#)
- *Topic **distribution**
 - eulermultinom, [18](#)
- *Topic **interface**
 - Csnippet, [12](#)
 - pomp constructor, [46](#)
- *Topic **models**
 - blowflies, [10](#)
 - dacca, [16](#)
 - Example pomp models, [20](#)
 - gompertz, [22](#)
 - ou2, [38](#)
 - pomp constructor, [46](#)
 - POMP simulation, [57](#)
 - pomp-package, [3](#)
 - Process model plugins, [67](#)
 - ricker, [71](#)
 - rw2, [71](#)
 - sir, [73](#)
- *Topic **optimize**
 - Iterated filtering, [22](#)
 - Iterated filtering 2, [27](#)
 - Nonlinear forecasting, [36](#)
 - Power spectrum computation and matching, [58](#)
 - Probes and synthetic likelihood, [64](#)
 - Simulated annealing, [72](#)
 - Trajectory matching, [74](#)
- *Topic **programming**
 - Csnippet, [12](#)
 - Low-level-interface, [31](#)
 - MCMC proposal distributions, [35](#)
 - pomp constructor, [46](#)
 - pomp methods, [54](#)
- *Topic **smooth**
 - B-splines, [7](#)
- *Topic **ts**
 - Approximate Bayesian computation, [4](#)
 - Bayesian sequential Monte Carlo, [8](#)
 - Iterated filtering, [22](#)
 - Iterated filtering 2, [27](#)
 - Nonlinear forecasting, [36](#)
 - Particle filter, [40](#)
 - Particle Markov Chain Monte Carlo, [43](#)
 - pomp constructor, [46](#)
 - pomp methods, [54](#)
 - POMP simulation, [57](#)
 - pomp-package, [3](#)
 - Power spectrum computation and matching, [58](#)
 - Probe functions, [61](#)
 - Probes and synthetic likelihood, [64](#)
 - Trajectory matching, [74](#)
 - [,abcList-method (Approximate Bayesian computation), [4](#)
 - [,mif2List-method (Iterated filtering 2), [27](#)
 - [,mifList-method (Iterated filtering), [22](#)
 - [,pmcmcList-method (Particle Markov Chain Monte Carlo), [43](#)
 - [-abcList (Approximate Bayesian computation), [4](#)
 - [-mif2List (Iterated filtering 2), [27](#)

- `[-mifList` (Iterated filtering), 22
- `[-pmcmcList` (Particle Markov Chain Monte Carlo), 43
- `$,bsmcd.pomp-method` (Bayesian sequential Monte Carlo), 8
- `$,nlfd.pomp-method` (Nonlinear forecasting), 36
- `$,pfilterd.pomp-method` (Particle filter), 40
- `$,probe.matched.pomp-method` (Probes and synthetic likelihood), 64
- `$,probed.pomp-method` (Probes and synthetic likelihood), 64
- `$,traj.matched.pomp-method` (Trajectory matching), 74
- `-$bsmcd.pomp` (Bayesian sequential Monte Carlo), 8
- `-$nlfd.pomp` (Nonlinear forecasting), 36
- `-$pfilterd.pomp` (Particle filter), 40
- `-$probe.matched.pomp` (Probes and synthetic likelihood), 64
- `-$probed.pomp` (Probes and synthetic likelihood), 64
- `-$traj.matched.pomp` (Trajectory matching), 74
- ABC (Approximate Bayesian computation), 4
- `abc`, 3, 36
- `abc` (Approximate Bayesian computation), 4
- `abc,abc-method` (Approximate Bayesian computation), 4
- `abc,pomp-method` (Approximate Bayesian computation), 4
- `abc,probed.pomp-method` (Approximate Bayesian computation), 4
- `abc-abc` (Approximate Bayesian computation), 4
- `abc-class` (Approximate Bayesian computation), 4
- `abc-methods` (Approximate Bayesian computation), 4
- `abc-pomp` (Approximate Bayesian computation), 4
- `abc-probed.pomp` (Approximate Bayesian computation), 4
- `abcList-class` (Approximate Bayesian computation), 4
- accumulator variables (pomp constructor), 46
- Approximate Bayesian computation, 4
- `as,pfilterd.pomp-method` (Particle filter), 40
- `as,pomp-method` (pomp methods), 54
- `as,probed.pomp-method` (Probes and synthetic likelihood), 64
- `as.data.frame.pfilterd.pomp` (Particle filter), 40
- `as.data.frame.pomp` (pomp methods), 54
- B-splines, 7
- `bake` (Utilities for reproducibility), 76
- `basic.probes` (Probe functions), 61
- Bayesian sequential Monte Carlo, 8
- `bbs`, 21
- `bbs(sir)`, 73
- `blowflies`, 10, 21
- `blowflies1` (blowflies), 10
- `blowflies2` (blowflies), 10
- `bsmc`, 3
- `bsmc` (Bayesian sequential Monte Carlo), 8
- `bsmc,pomp-method` (Bayesian sequential Monte Carlo), 8
- `bsmc-pomp` (Bayesian sequential Monte Carlo), 8
- `bsmc2`, 4, 43
- `bsmc2` (Bayesian sequential Monte Carlo), 8
- `bsmc2,pomp-method` (Bayesian sequential Monte Carlo), 8
- `bsmc2-pomp` (Bayesian sequential Monte Carlo), 8
- `bspline.basis` (B-splines), 7
- `c,abc-method` (Approximate Bayesian computation), 4
- `c,abcList-method` (Approximate Bayesian computation), 4
- `c,mif-method` (Iterated filtering), 22
- `c,mif2d.pomp-method` (Iterated filtering 2), 27
- `c,mif2List-method` (Iterated filtering 2), 27
- `c,mifList-method` (Iterated filtering), 22

- c,pmcmc-method (Particle Markov Chain Monte Carlo), [43](#)
- c,pmcmcList-method (Particle Markov Chain Monte Carlo), [43](#)
- c-abc (Approximate Bayesian computation), [4](#)
- c-abcList (Approximate Bayesian computation), [4](#)
- c-mif (Iterated filtering), [22](#)
- c-mif2d.pomp (Iterated filtering 2), [27](#)
- c-mif2List (Iterated filtering 2), [27](#)
- c-mifList (Iterated filtering), [22](#)
- c-pmcmc (Particle Markov Chain Monte Carlo), [43](#)
- c-pmcmcList (Particle Markov Chain Monte Carlo), [43](#)
- Childhood disease incidence data, [11](#)
- coef,mif2List-method (Iterated filtering 2), [27](#)
- coef,mifList-method (Iterated filtering), [22](#)
- coef,pomp-method (pomp methods), [54](#)
- coef-pomp (pomp methods), [54](#)
- coef.rec-mif2List (Iterated filtering 2), [27](#)
- coef.rec-mifList (Iterated filtering), [22](#)
- coef<- (pomp methods), [54](#)
- coef<-,pomp-method (pomp methods), [54](#)
- coef<--pomp (pomp methods), [54](#)
- coerce,pfilterd.pomp,data.frame-method (Particle filter), [40](#)
- coerce,pomp,data.frame-method (pomp methods), [54](#)
- coerce,probed.pomp,data.frame-method (Probes and synthetic likelihood), [64](#)
- cond.logLik (Particle filter), [40](#)
- cond.logLik,pfilterd.pomp-method (Particle filter), [40](#)
- cond.logLik-pfilterd.pomp (Particle filter), [40](#)
- continue (Iterated filtering), [22](#)
- continue,abc-method (Approximate Bayesian computation), [4](#)
- continue,mif-method (Iterated filtering), [22](#)
- continue,mif2d.pomp-method (Iterated filtering 2), [27](#)
- continue,pmcmc-method (Particle Markov Chain Monte Carlo), [43](#)
- continue-abc (Approximate Bayesian computation), [4](#)
- continue-mif (Iterated filtering), [22](#)
- continue-mif2d.pomp (Iterated filtering 2), [27](#)
- continue-pmcmc (Particle Markov Chain Monte Carlo), [43](#)
- conv.rec (Iterated filtering), [22](#)
- conv.rec,abc-method (Approximate Bayesian computation), [4](#)
- conv.rec,abcList-method (Approximate Bayesian computation), [4](#)
- conv.rec,mif-method (Iterated filtering), [22](#)
- conv.rec,mif2d.pomp-method (Iterated filtering 2), [27](#)
- conv.rec,mif2List-method (Iterated filtering 2), [27](#)
- conv.rec,mifList-method (Iterated filtering), [22](#)
- conv.rec,pmcmc-method (Particle Markov Chain Monte Carlo), [43](#)
- conv.rec,pmcmcList-method (Particle Markov Chain Monte Carlo), [43](#)
- conv.rec-abc (Approximate Bayesian computation), [4](#)
- conv.rec-abcList (Approximate Bayesian computation), [4](#)
- conv.rec-mif (Iterated filtering), [22](#)
- conv.rec-mif2d.pomp (Iterated filtering 2), [27](#)
- conv.rec-mif2List (Iterated filtering 2), [27](#)
- conv.rec-mifList (Iterated filtering), [22](#)
- conv.rec-pmcmc (Particle Markov Chain Monte Carlo), [43](#)
- conv.rec-pmcmcList (Particle Markov Chain Monte Carlo), [43](#)
- covmat (Approximate Bayesian computation), [4](#)
- covmat,abc-method (Approximate Bayesian computation), [4](#)
- covmat,abcList-method (Approximate Bayesian computation), [4](#)

- covmat,pmcmc-method (Particle Markov Chain Monte Carlo), 43
- covmat,pmcmcList-method (Particle Markov Chain Monte Carlo), 43
- covmat-abc (Approximate Bayesian computation), 4
- covmat-abcList (Approximate Bayesian computation), 4
- covmat-pmcmc (Particle Markov Chain Monte Carlo), 43
- covmat-pmcmcList (Particle Markov Chain Monte Carlo), 43
- Csnippet, 12, 46, 48–50, 52, 68, 69
- Csnippet-class (Csnippet), 12

- dacca, 16, 21
- data.frame-pomp (pomp constructor), 46
- design, 17
- deSolve, 34
- deulermultinom (eulermultinom), 18
- discrete.time.sim, 51
- discrete.time.sim (Process model plugins), 67
- dmeasure (Low-level-interface), 31
- dmeasure,pomp-method (Low-level-interface), 31
- dmeasure-pomp (Low-level-interface), 31
- dprior (Low-level-interface), 31
- dprior,pomp-method (Low-level-interface), 31
- dprior-pomp (Low-level-interface), 31
- dprocess (Low-level-interface), 31
- dprocess,pomp-method (Low-level-interface), 31
- dprocess-pomp (Low-level-interface), 31

- eff.sample.size (Particle filter), 40
- eff.sample.size,pfilterd.pomp-method (Particle filter), 40
- eff.sample.size-pfilterd.pomp (Particle filter), 40
- euler.sim, 51
- euler.sim (Process model plugins), 67
- euler.sir, 16, 21
- euler.sir (sir), 73
- eulermultinom, 18
- ewcitmeas (Childhood disease incidence data), 11
- ewmeas (Childhood disease incidence data), 11
- Example pomp models, 20

- filter.mean (Particle filter), 40
- filter.mean,pfilterd.pomp-method (Particle filter), 40
- filter.mean-pfilterd.pomp (Particle filter), 40
- filter.traj (Particle filter), 40
- filter.traj,pfilterd.pomp-method (Particle filter), 40
- filter.traj-pfilterd.pomp (Particle filter), 40
- freeze (Utilities for reproducibility), 76

- gillespie.sim, 51
- gillespie.sim (Process model plugins), 67
- gillespie.sir, 21
- gillespie.sir (sir), 73
- gompertz, 21, 22, 71

- here, 3

- init.state (Low-level-interface), 31
- init.state,pomp-method (Low-level-interface), 31
- init.state-pomp (Low-level-interface), 31
- Iterated filtering, 22
- Iterated filtering 2, 27

- kernel, 62

- load, 77
- logLik,nlfd.pomp-method (Nonlinear forecasting), 36
- logLik,pfilterd.pomp-method (Particle filter), 40
- logLik,pmcmc-method (Particle Markov Chain Monte Carlo), 43
- logLik,probed.pomp-method (Probes and synthetic likelihood), 64
- logLik,traj.matched.pomp-method (Trajectory matching), 74
- logLik-nlfd.pomp (Nonlinear forecasting), 36

- logLik-pfilterd.pomp (Particle filter), 40
- logLik-pmcmc (Particle Markov Chain Monte Carlo), 43
- logLik-probed.pomp (Probes and synthetic likelihood), 64
- logLik-traj.matched.pomp (Trajectory matching), 74
- logmeanexp, 30
- LondonYorke (Childhood disease incidence data), 11
- Low-level-interface, 31

- matrix-pomp (pomp constructor), 46
- mcmc, 45
- MCMC proposal distributions, 6, 35, 45
- MCMC proposal functions, 5, 44
- MCMC proposal functions (MCMC proposal distributions), 35
- mcmc.list, 45
- mean, 62
- mif, 3, 4, 28, 30, 43
- mif (Iterated filtering), 22
- mif, mif-method (Iterated filtering), 22
- mif, pfilterd.pomp-method (Iterated filtering), 22
- mif, pomp-method (Iterated filtering), 22
- mif-class (Iterated filtering), 22
- mif-methods (Iterated filtering), 22
- mif-mif (Iterated filtering), 22
- mif-pfilterd.pomp (Iterated filtering), 22
- mif-pomp (Iterated filtering), 22
- mif2, 3, 22, 25, 26
- mif2 (Iterated filtering 2), 27
- mif2, mif2d.pomp-method (Iterated filtering 2), 27
- mif2, pfilterd.pomp-method (Iterated filtering 2), 27
- mif2, pomp-method (Iterated filtering 2), 27
- mif2-mif2d.pomp (Iterated filtering 2), 27
- mif2-pfilterd.pomp (Iterated filtering 2), 27
- mif2-pomp (Iterated filtering 2), 27
- mif2d.pomp-class (Iterated filtering 2), 27
- mif2d.pomp-methods (Iterated filtering 2), 27
- mif2List-class (Iterated filtering 2), 27
- mifList-class (Iterated filtering), 22
- mvn.diag.rw (MCMC proposal distributions), 35
- mvn.rw, 6, 45
- mvn.rw (MCMC proposal distributions), 35

- nlf, 3, 4
- nlf (Nonlinear forecasting), 36
- nlf, nlf.d.pomp-method (Nonlinear forecasting), 36
- nlf, pomp-method (Nonlinear forecasting), 36
- nlf-nlf.d.pomp (Nonlinear forecasting), 36
- nlf-pomp (Nonlinear forecasting), 36
- nlf.d.pomp-class (Nonlinear forecasting), 36
- nloptr, 65, 66, 75
- Nonlinear forecasting, 36
- numeric-pomp (pomp constructor), 46

- obs, 63
- obs (pomp methods), 54
- obs, pomp-method (pomp methods), 54
- obs-pomp (pomp methods), 54
- ode, 33, 34
- onestep.dens, 51
- onestep.dens (Process model plugins), 67
- onestep.sim, 51
- onestep.sim (Process model plugins), 67
- optim, 37, 60, 65, 66, 73, 75, 76
- ou2, 21, 38, 71

- parmat, 39
- Particle filter, 40
- particle filter (Particle filter), 40
- Particle Markov Chain Monte Carlo, 43
- partrans, 34
- partrans (pomp methods), 54
- partrans, pomp-method (pomp methods), 54
- partrans-pomp (pomp methods), 54
- paste, 7
- periodic.bspline.basis (B-splines), 7
- pfilter, 3, 4, 10, 24, 26, 28, 30, 45, 57
- pfilter (Particle filter), 40

- pfilter, pfilterd.pomp-method (Particle filter), [40](#)
- pfilter, pomp-method (Particle filter), [40](#)
- pfilter-pfilterd.pomp (Particle filter), [40](#)
- pfilter-pomp (Particle filter), [40](#)
- pfilterd.pomp, [25](#), [28](#), [41](#)
- pfilterd.pomp (Particle filter), [40](#)
- pfilterd.pomp-class (Particle filter), [40](#)
- plot, abc-method (Approximate Bayesian computation), [4](#)
- plot, abcList-method (Approximate Bayesian computation), [4](#)
- plot, bsmcd.pomp-method (Bayesian sequential Monte Carlo), [8](#)
- plot, mif-method (Iterated filtering), [22](#)
- plot, mif2d.pomp-method (Iterated filtering 2), [27](#)
- plot, mif2List-method (Iterated filtering 2), [27](#)
- plot, mifList-method (Iterated filtering), [22](#)
- plot, pmcmc-method (Particle Markov Chain Monte Carlo), [43](#)
- plot, pmcmcList-method (Particle Markov Chain Monte Carlo), [43](#)
- plot, pomp-method (pomp methods), [54](#)
- plot, probe.matched.pomp-method (Probes and synthetic likelihood), [64](#)
- plot, probed.pomp-method (Probes and synthetic likelihood), [64](#)
- plot, spect.matched.pomp-method (Probes and synthetic likelihood), [64](#)
- plot, spect.pomp-method (Probes and synthetic likelihood), [64](#)
- plot-abc (Approximate Bayesian computation), [4](#)
- plot-abcList (Approximate Bayesian computation), [4](#)
- plot-bsmcd.pomp (Bayesian sequential Monte Carlo), [8](#)
- plot-mif (Iterated filtering), [22](#)
- plot-mif2d.pomp (Iterated filtering 2), [27](#)
- plot-mif2List (Iterated filtering 2), [27](#)
- plot-mifList (Iterated filtering), [22](#)
- plot-pmcmc (Particle Markov Chain Monte Carlo), [43](#)
- plot-pmcmcList (Particle Markov Chain Monte Carlo), [43](#)
- plot-pomp (pomp methods), [54](#)
- plot-probe.matched.pomp (Probes and synthetic likelihood), [64](#)
- plot-probed.pomp (Probes and synthetic likelihood), [64](#)
- plot-spect.pomp (Probes and synthetic likelihood), [64](#)
- plugins, [15](#), [47](#), [51](#)
- plugins (Process model plugins), [67](#)
- pmcmc, [3](#), [4](#), [36](#), [43](#)
- pmcmc (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc, pfilterd.pomp-method (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc, pmcmc-method (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc, pomp-method (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc-class (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc-methods (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc-pfilterd.pomp (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc-pmcmc (Particle Markov Chain Monte Carlo), [43](#)
- pmcmc-pomp (Particle Markov Chain Monte Carlo), [43](#)
- pmcmcList-class (Particle Markov Chain Monte Carlo), [43](#)
- pomp, [3](#), [4](#), [6](#), [10](#), [11](#), [13](#), [15](#), [16](#), [24–26](#), [28](#), [30](#), [35](#), [39](#), [41](#), [43](#), [45](#), [51](#), [57](#), [58](#), [60](#), [61](#), [63](#), [65](#), [66](#), [70](#), [71](#), [73–76](#)
- pomp (pomp constructor), [46](#)
- pomp constructor, [46](#)
- pomp low-level interface, [4](#), [53](#), [57](#)
- pomp low-level interface (Low-level-interface), [31](#)
- pomp methods, [35](#), [53](#), [54](#)
- pomp package (pomp-package), [3](#)
- POMP simulation, [57](#)
- pomp, data.frame-method (pomp constructor), [46](#)
- pomp, matrix-method (pomp constructor),

- 46
- pomp,numeric-method (pomp constructor), 46
- pomp,pomp-method (pomp constructor), 46
- pomp-class (pomp constructor), 46
- pomp-methods (pomp methods), 54
- pomp-package, 3
- pomp-pomp (pomp constructor), 46
- pompExample, 3, 50
- pompExample (Example pomp models), 20
- pompLoad (Low-level-interface), 31
- pompLoad,pomp-method (Low-level-interface), 31
- pompLoad-pomp (Low-level-interface), 31
- pompUnload (Low-level-interface), 31
- pompUnload,pomp-method (Low-level-interface), 31
- pompUnload-pomp (Low-level-interface), 31
- Power spectrum computation and matching, 58
- power spectrum computation and matching (Power spectrum computation and matching), 58
- pred.mean (Particle filter), 40
- pred.mean,pfilterd.pomp-method (Particle filter), 40
- pred.mean-pfilterd.pomp (Particle filter), 40
- pred.var (Particle filter), 40
- pred.var,pfilterd.pomp-method (Particle filter), 40
- pred.var-pfilterd.pomp (Particle filter), 40
- print,pomp-method (pomp methods), 54
- print-pomp (pomp methods), 54
- probe, 3–6, 57, 61, 63
- probe (Probes and synthetic likelihood), 64
- Probe functions, 61
- probe functions, 65, 66
- probe functions (Probe functions), 61
- probe,pomp-method (Probes and synthetic likelihood), 64
- probe,probed.pomp-method (Probes and synthetic likelihood), 64
- probe-pomp (Probes and synthetic likelihood), 64
- probe-probed.pomp (Probes and synthetic likelihood), 64
- probe.acf (Probe functions), 61
- probe.ccf (Probe functions), 61
- probe.marginal (Probe functions), 61
- probe.match, 3, 63, 73
- probe.match (Probes and synthetic likelihood), 64
- probe.match,pomp-method (Probes and synthetic likelihood), 64
- probe.match,probe.matched.pomp-method (Probes and synthetic likelihood), 64
- probe.match,probed.pomp-method (Probes and synthetic likelihood), 64
- probe.match-pomp (Probes and synthetic likelihood), 64
- probe.match-probe.matched.pomp (Probes and synthetic likelihood), 64
- probe.match-probed.pomp (Probes and synthetic likelihood), 64
- probe.match.objfun (Probes and synthetic likelihood), 64
- probe.match.objfun,pomp-method (Probes and synthetic likelihood), 64
- probe.match.objfun,probed.pomp-method (Probes and synthetic likelihood), 64
- probe.match.objfun-pomp (Probes and synthetic likelihood), 64
- probe.match.objfun-probed.pomp (Probes and synthetic likelihood), 64
- probe.matched.pomp, 65
- probe.matched.pomp-class (Probes and synthetic likelihood), 64
- probe.matched.pomp-methods (Probes and synthetic likelihood), 64
- probe.mean (Probe functions), 61
- probe.median (Probe functions), 61
- probe.nlar (Probe functions), 61
- probe.period (Probe functions), 61
- probe.quantile (Probe functions), 61
- probe.sd (Probe functions), 61
- probe.var (Probe functions), 61
- probed.pomp-class (Probes and synthetic likelihood), 64
- probed.pomp-methods (Probes and synthetic likelihood), 64

- Probes and synthetic likelihood, [64](#)
- Process model plugins, [67](#)
- process model plugins, [53](#)
- process model plugins (Process model plugins), [67](#)
- profileDesign (design), [17](#)
- quantile, [62](#)
- readRDS, [77](#)
- reulermultinom (eulermultinom), [18](#)
- rgammawn (eulermultinom), [18](#)
- ricker, [21](#), [71](#)
- rmeasure (Low-level-interface), [31](#)
- rmeasure, pomp-method (Low-level-interface), [31](#)
- rmeasure-pomp (Low-level-interface), [31](#)
- rprior, [9](#)
- rprior (Low-level-interface), [31](#)
- rprior, pomp-method (Low-level-interface), [31](#)
- rprior-pomp (Low-level-interface), [31](#)
- rprocess (Low-level-interface), [31](#)
- rprocess, pomp-method (Low-level-interface), [31](#)
- rprocess-pomp (Low-level-interface), [31](#)
- rw.sd (Iterated filtering 2), [27](#)
- rw2, [21](#), [71](#)
- sannbox, [65](#), [75](#)
- sannbox (Simulated annealing), [72](#)
- sequential Monte Carlo (Particle filter), [40](#)
- set.seed, [77](#)
- show, pomp-method (pomp methods), [54](#)
- show-pomp (pomp methods), [54](#)
- simulate, [3](#), [4](#), [57](#), [59](#), [65](#)
- simulate, pomp-method (POMP simulation), [57](#)
- simulate-pomp, [65](#)
- simulate-pomp (POMP simulation), [57](#)
- Simulated annealing, [72](#)
- sir, [73](#)
- skeleton (Low-level-interface), [31](#)
- skeleton, pomp-method (Low-level-interface), [31](#)
- skeleton-pomp (Low-level-interface), [31](#)
- sliceDesign (design), [17](#)
- SMC (Particle filter), [40](#)
- sobol (design), [17](#)
- sobolDesign (design), [17](#)
- spect, [66](#)
- spect (Power spectrum computation and matching), [58](#)
- spect, pomp-method (Power spectrum computation and matching), [58](#)
- spect, spect.pomp-method (Power spectrum computation and matching), [58](#)
- spect-pomp (Power spectrum computation and matching), [58](#)
- spect-spect.pomp (Power spectrum computation and matching), [58](#)
- spect.match, [3](#)
- spect.match (Power spectrum computation and matching), [58](#)
- spect.match, pomp-method (Power spectrum computation and matching), [58](#)
- spect.match, spect.pomp-method (Power spectrum computation and matching), [58](#)
- spect.match-pomp (Power spectrum computation and matching), [58](#)
- spect.match-spect.pomp (Power spectrum computation and matching), [58](#)
- spect.matched.pomp, [60](#)
- spect.matched.pomp-class (Power spectrum computation and matching), [58](#)
- spect.matched.pomp-methods (Probes and synthetic likelihood), [64](#)
- spect.pomp, [60](#)
- spect.pomp-class (Power spectrum computation and matching), [58](#)
- spect.pomp-methods (Probes and synthetic likelihood), [64](#)
- sprintf, [7](#)
- states (pomp methods), [54](#)
- states, pomp-method (pomp methods), [54](#)
- states-pomp (pomp methods), [54](#)
- stew (Utilities for reproducibility), [76](#)
- subplex, [37](#), [60](#), [65](#), [75](#), [76](#)
- summary, probe.matched.pomp-method (Probes and synthetic likelihood), [64](#)
- summary, probed.pomp-method (Probes and

- synthetic likelihood), 64
- summary, spect.matched.pomp-method (Probes and synthetic likelihood), 64
- summary, spect.pomp-method (Probes and synthetic likelihood), 64
- summary, traj.matched.pomp-method (Trajectory matching), 74
- summary-probe.matched.pomp (Probes and synthetic likelihood), 64
- summary-probed.pomp (Probes and synthetic likelihood), 64
- summary-spect.matched.pomp (Probes and synthetic likelihood), 64
- summary-spect.pomp (Probes and synthetic likelihood), 64
- summary-traj.matched.pomp (Trajectory matching), 74
- The pomp package (pomp-package), 3
- time, pomp-method (pomp methods), 54
- time-pomp (pomp methods), 54
- time<- (pomp methods), 54
- time<-, pomp-method (pomp methods), 54
- time<--pomp (pomp methods), 54
- timezero (pomp methods), 54
- timezero, pomp-method (pomp methods), 54
- timezero-pomp (pomp methods), 54
- timezero<- (pomp methods), 54
- timezero<-, pomp-method (pomp methods), 54
- timezero<--pomp (pomp methods), 54
- traj.match, 3, 4, 73
- traj.match (Trajectory matching), 74
- traj.match, pomp-method (Trajectory matching), 74
- traj.match, traj.matched.pomp-method (Trajectory matching), 74
- traj.match-pomp (Trajectory matching), 74
- traj.match-traj.matched.pomp (Trajectory matching), 74
- traj.match.objfun (Trajectory matching), 74
- traj.match.objfun, pomp-method (Trajectory matching), 74
- traj.match.objfun-pomp (Trajectory matching), 74
- traj.matched.pomp-class (Trajectory matching), 74
- trajectory, 75, 76
- trajectory (Low-level-interface), 31
- Trajectory matching, 74
- trajectory, pomp-method (Low-level-interface), 31
- trajectory-pomp (Low-level-interface), 31
- Utilities for reproducibility, 76
- values (Probes and synthetic likelihood), 64
- values, probe.matched.pomp-method (Probes and synthetic likelihood), 64
- values, probed.pomp-method (Probes and synthetic likelihood), 64
- values-probe.matched.pomp (Probes and synthetic likelihood), 64
- values-probed.pomp (Probes and synthetic likelihood), 64
- window, pomp-method (pomp methods), 54
- window-pomp (pomp methods), 54